

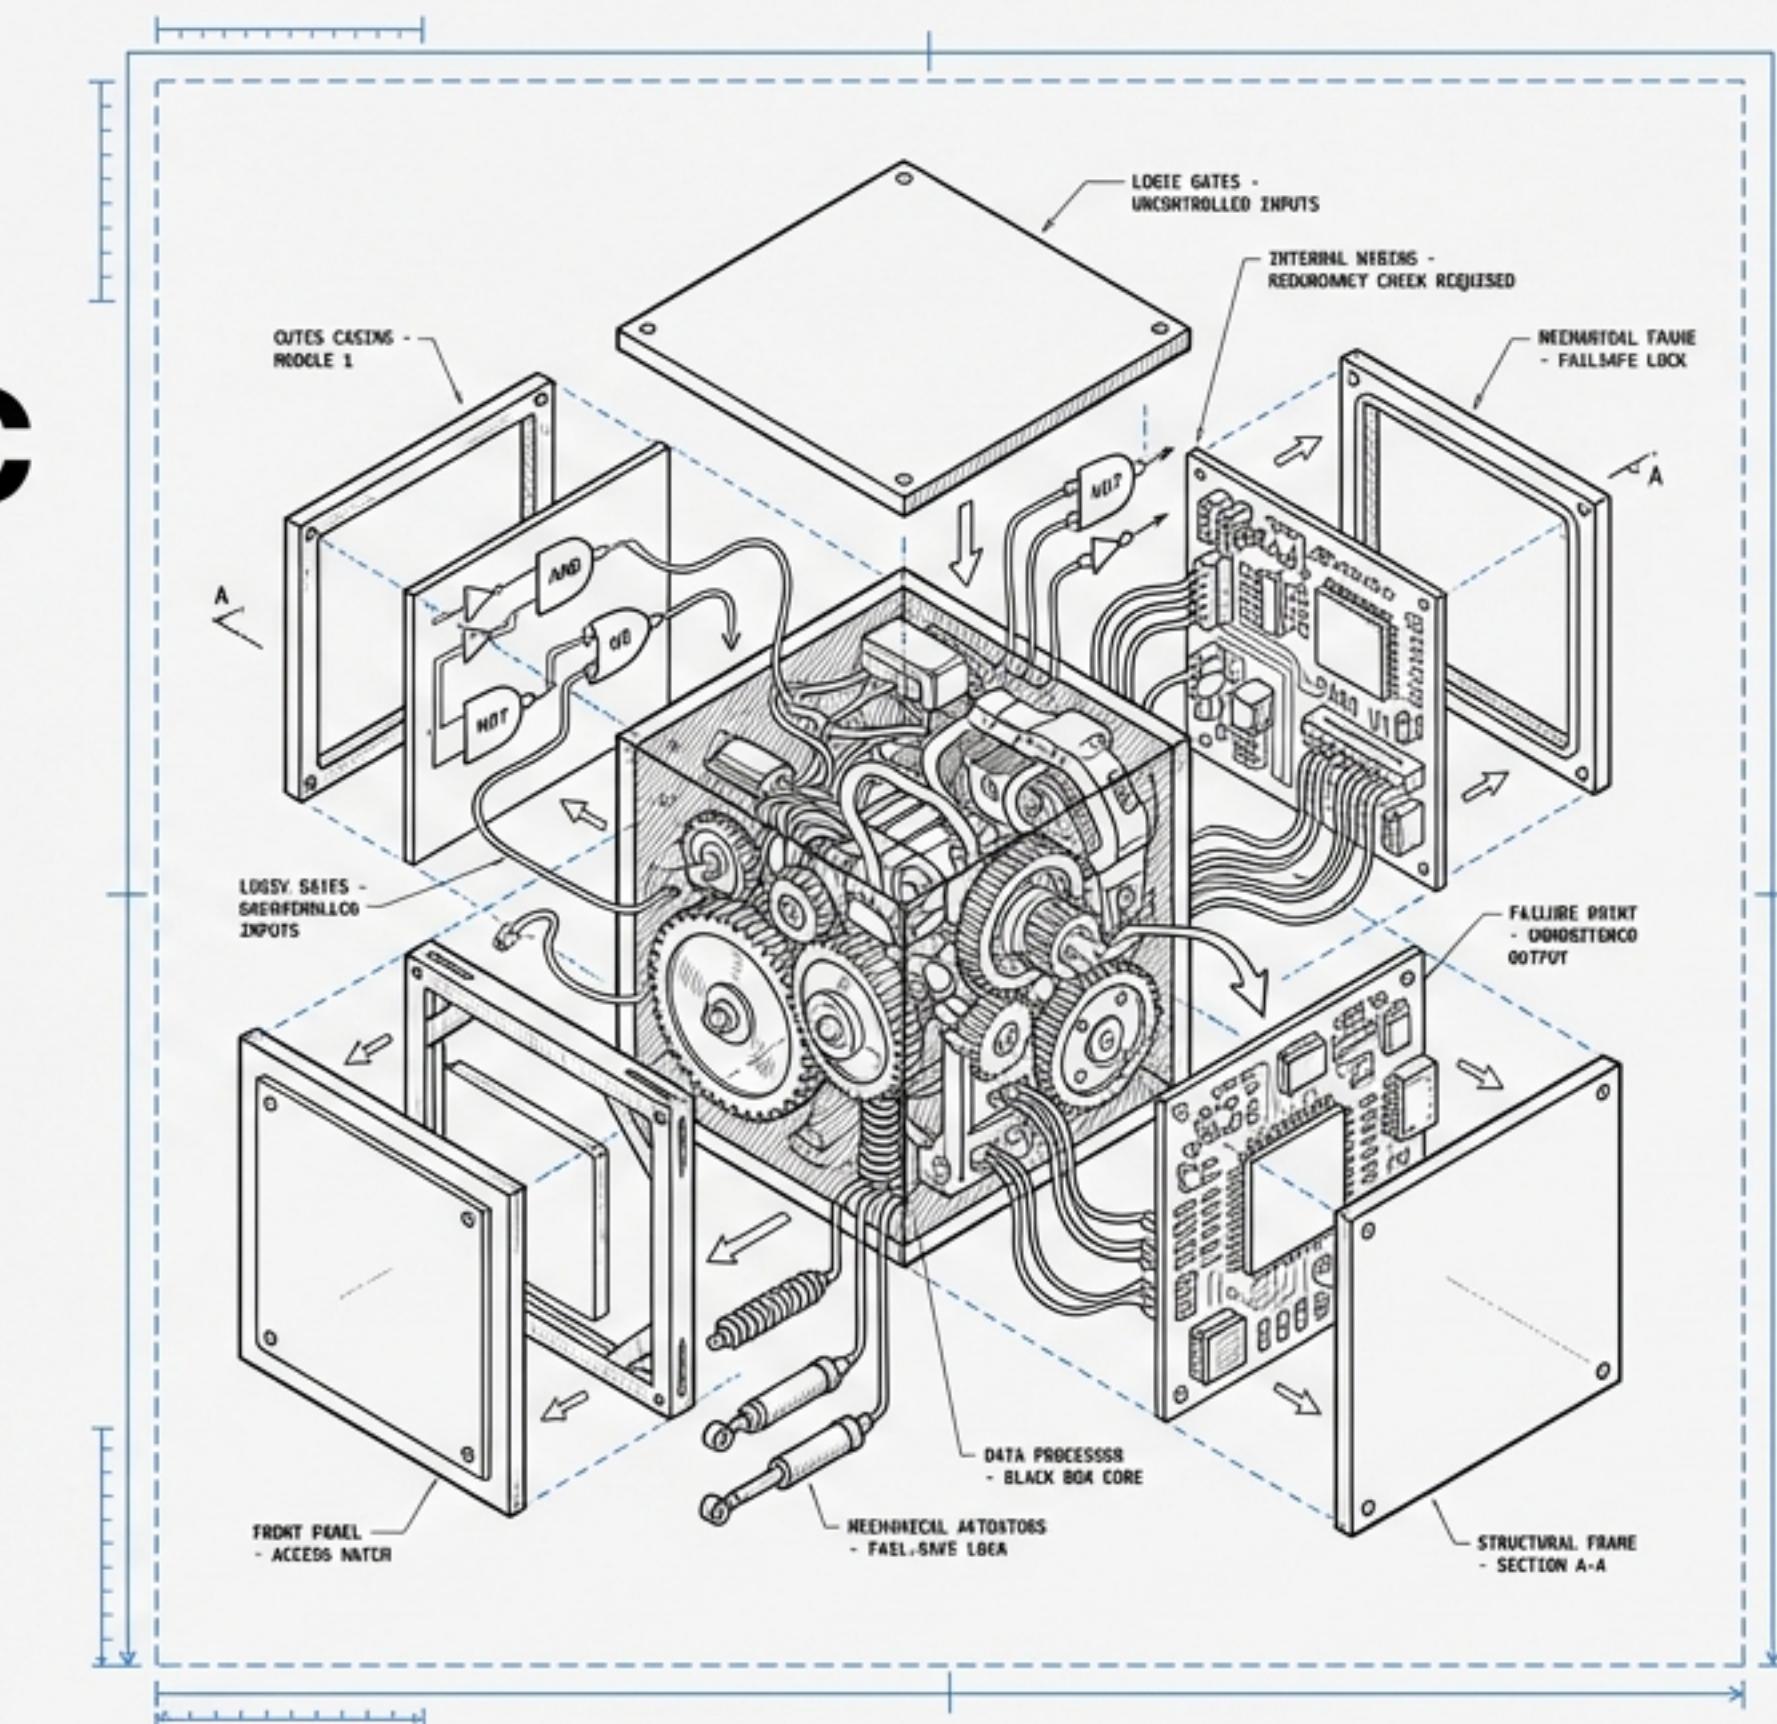
# AGENTIC AI: FROM BUZZWORDS TO ENGINEERING LOGIC

A Framework for Governance, Safety, and Organisational Maturity

## CONTEXT SPECIFICATION

**PURPOSE:** To decommission executive hype and replace it with structural constraints.

**CORE ARGUMENT:** Maturity is not measured by autonomy, but by how deliberately failure is engineered.



# A HARD BOUNDARY: WHAT ACTUALLY QUALIFIES AS AN AGENT?

## 5 System Properties



**Goal-Oriented:** Works toward an explicit objective.



**Planning Capability:** Breaks goals into steps and revises plans.



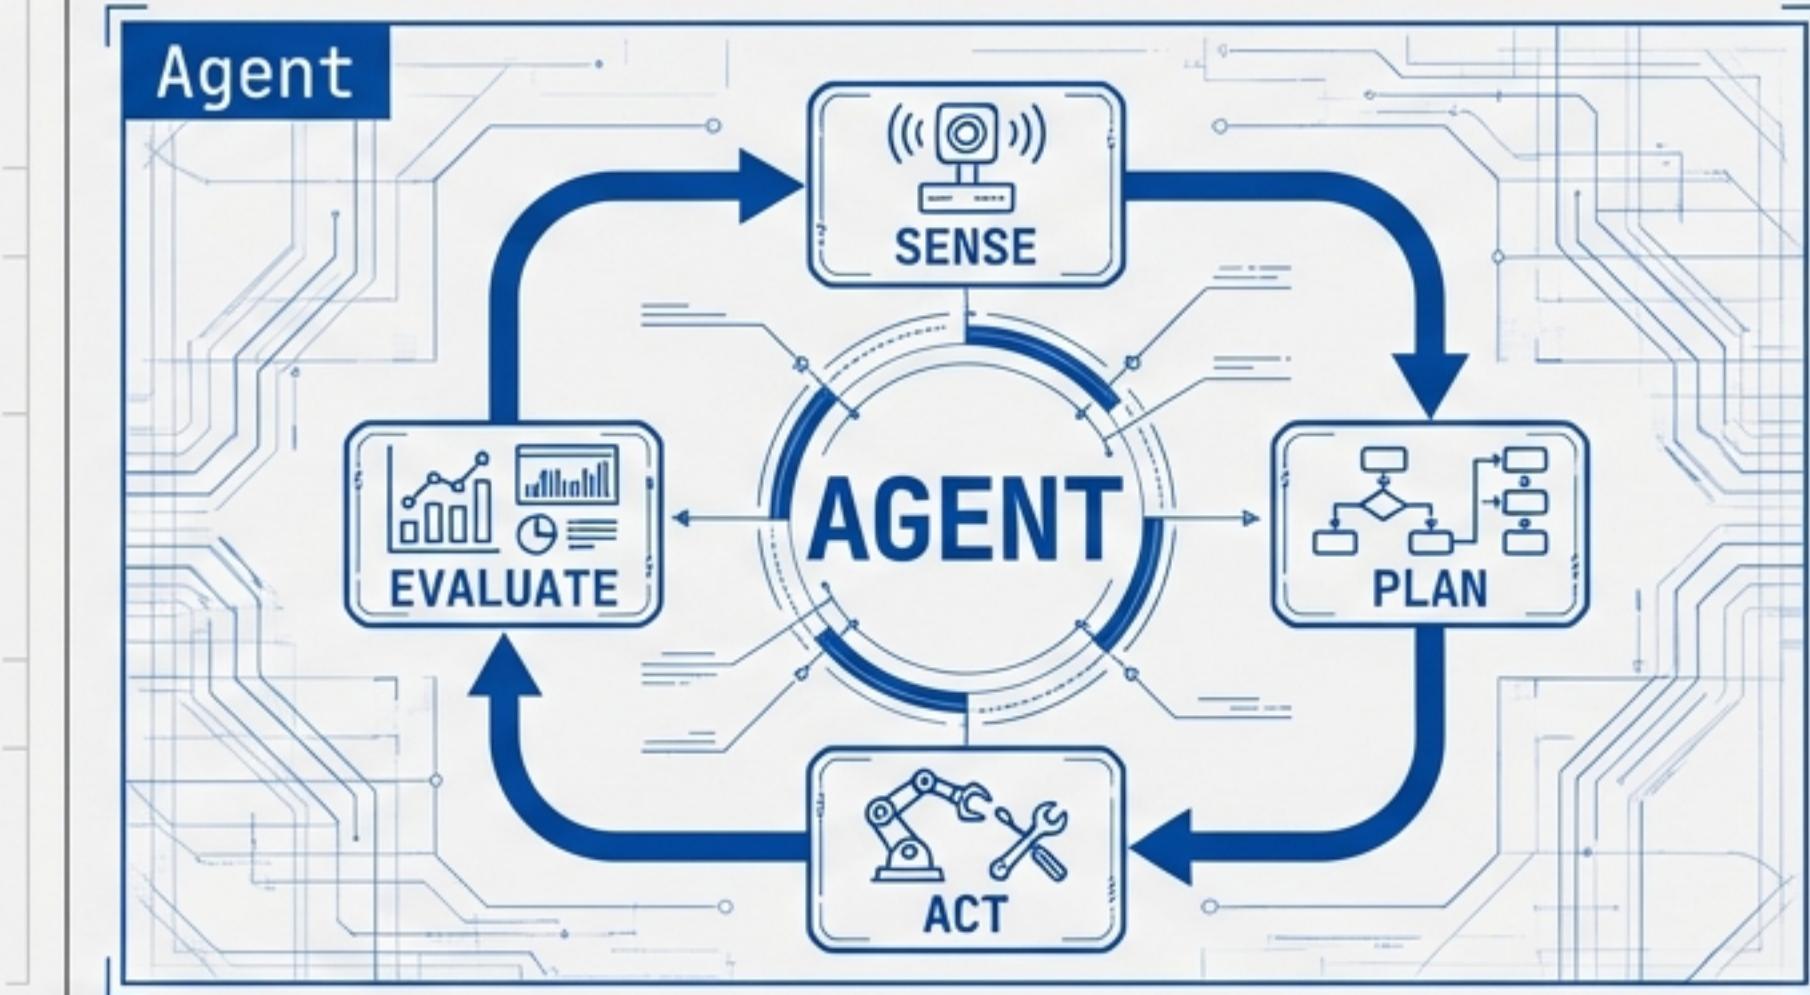
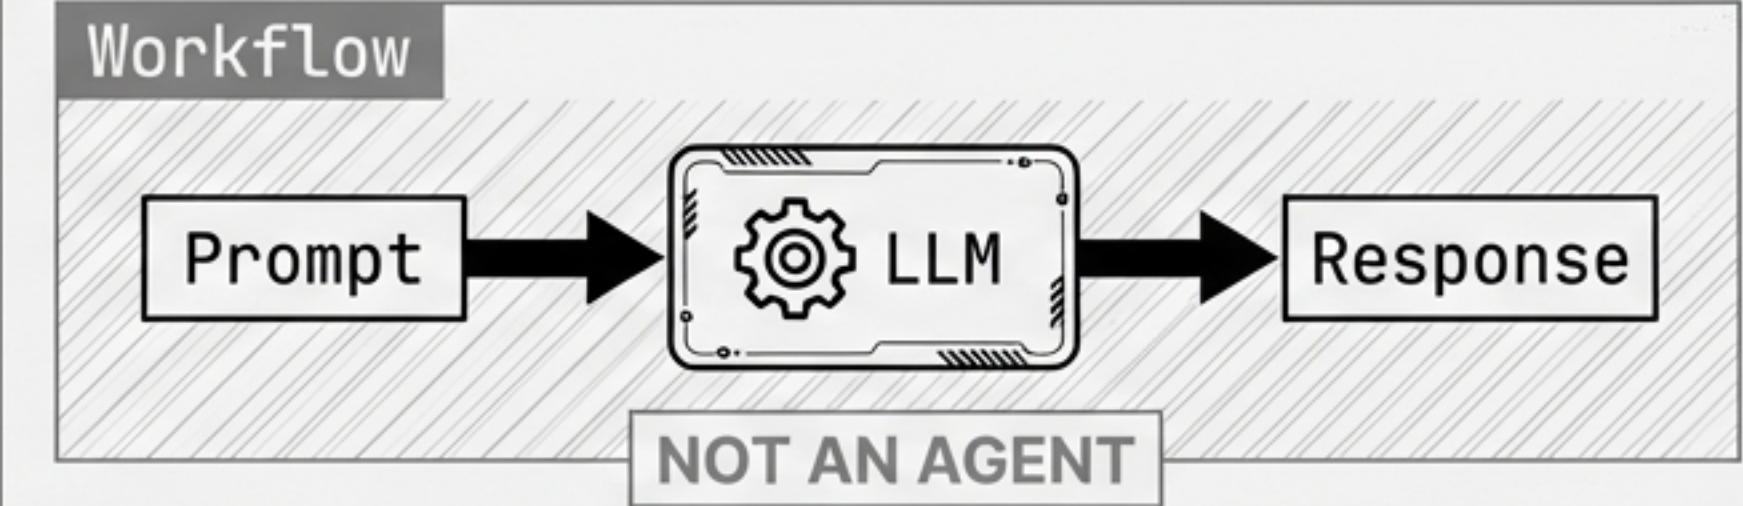
**Action Capability:** Executes API calls and tool usage.



**State & Memory:** Retains context across steps.



**Closed Feedback Loop:**  
Sense → Act → Evaluate → Adjust.



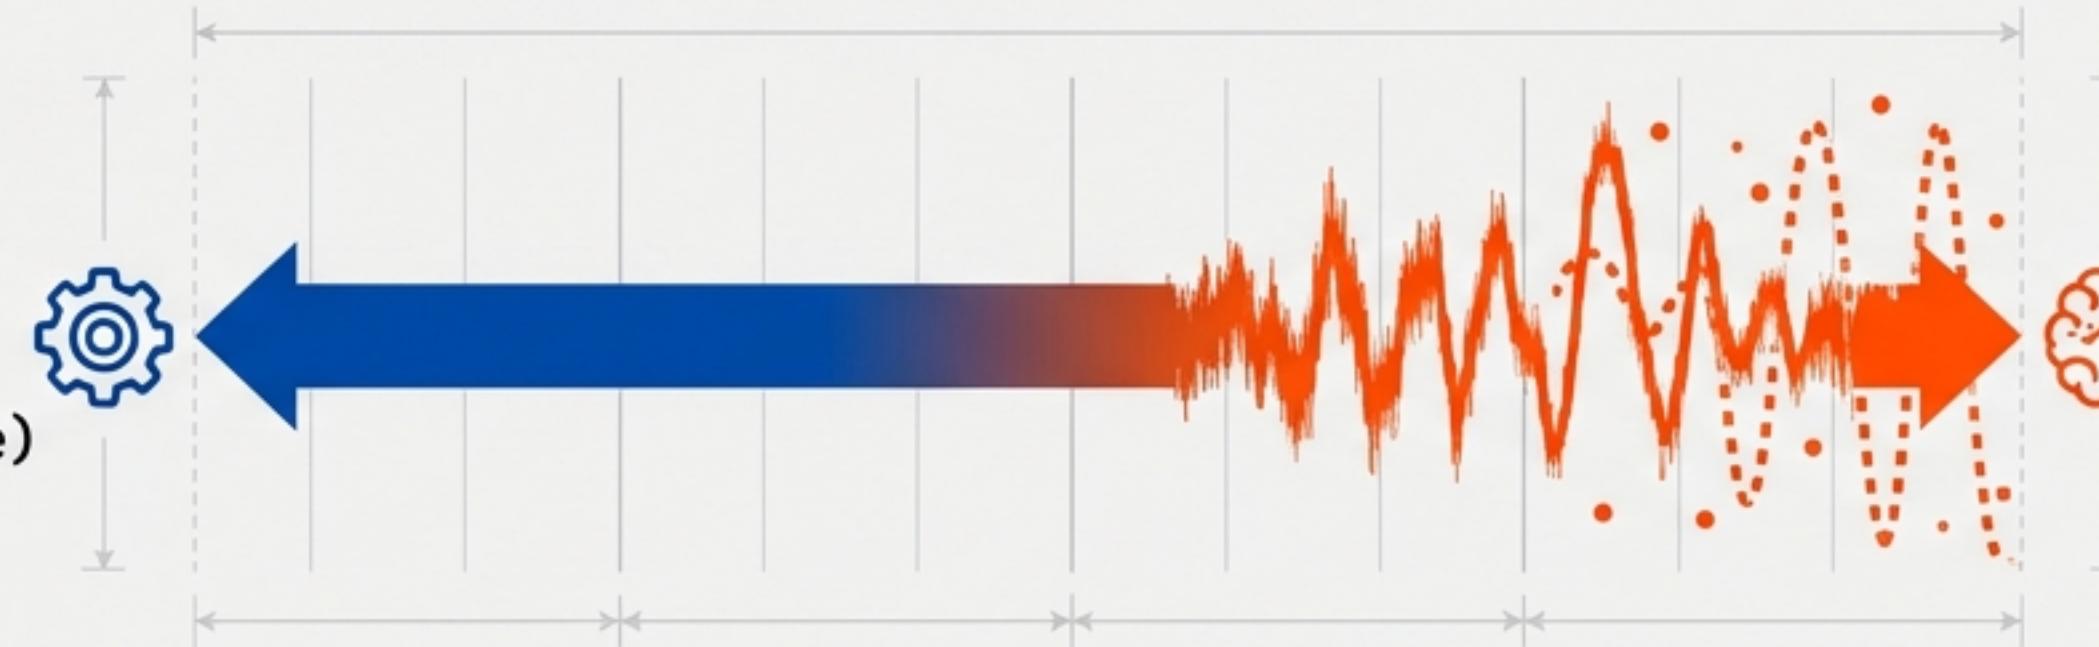
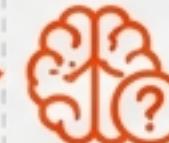
PROJECT:	AGENT QUALIFICATION CRITERIA	
CREATED BY:	SYSTEM ARCHITECT	REV: 1.0
DRAFTS:		SHEET: 1
DATE:	2024-10-26	

# THE DANGEROUS ILLUSION OF THE 'AUTONOMOUS WORKER'

**DETERMINISTIC SYSTEMS**  
(Traditional Software)



**PROBABILISTIC SYSTEMS**  
(Agentic AI)



Predictability	<input checked="" type="checkbox"/> 100% Consistent	Bounded & Variable
Explainability	High (Code-based)	Low (Model-based)
Cost Model	Linear / Stable	Non-Linear / Unstable



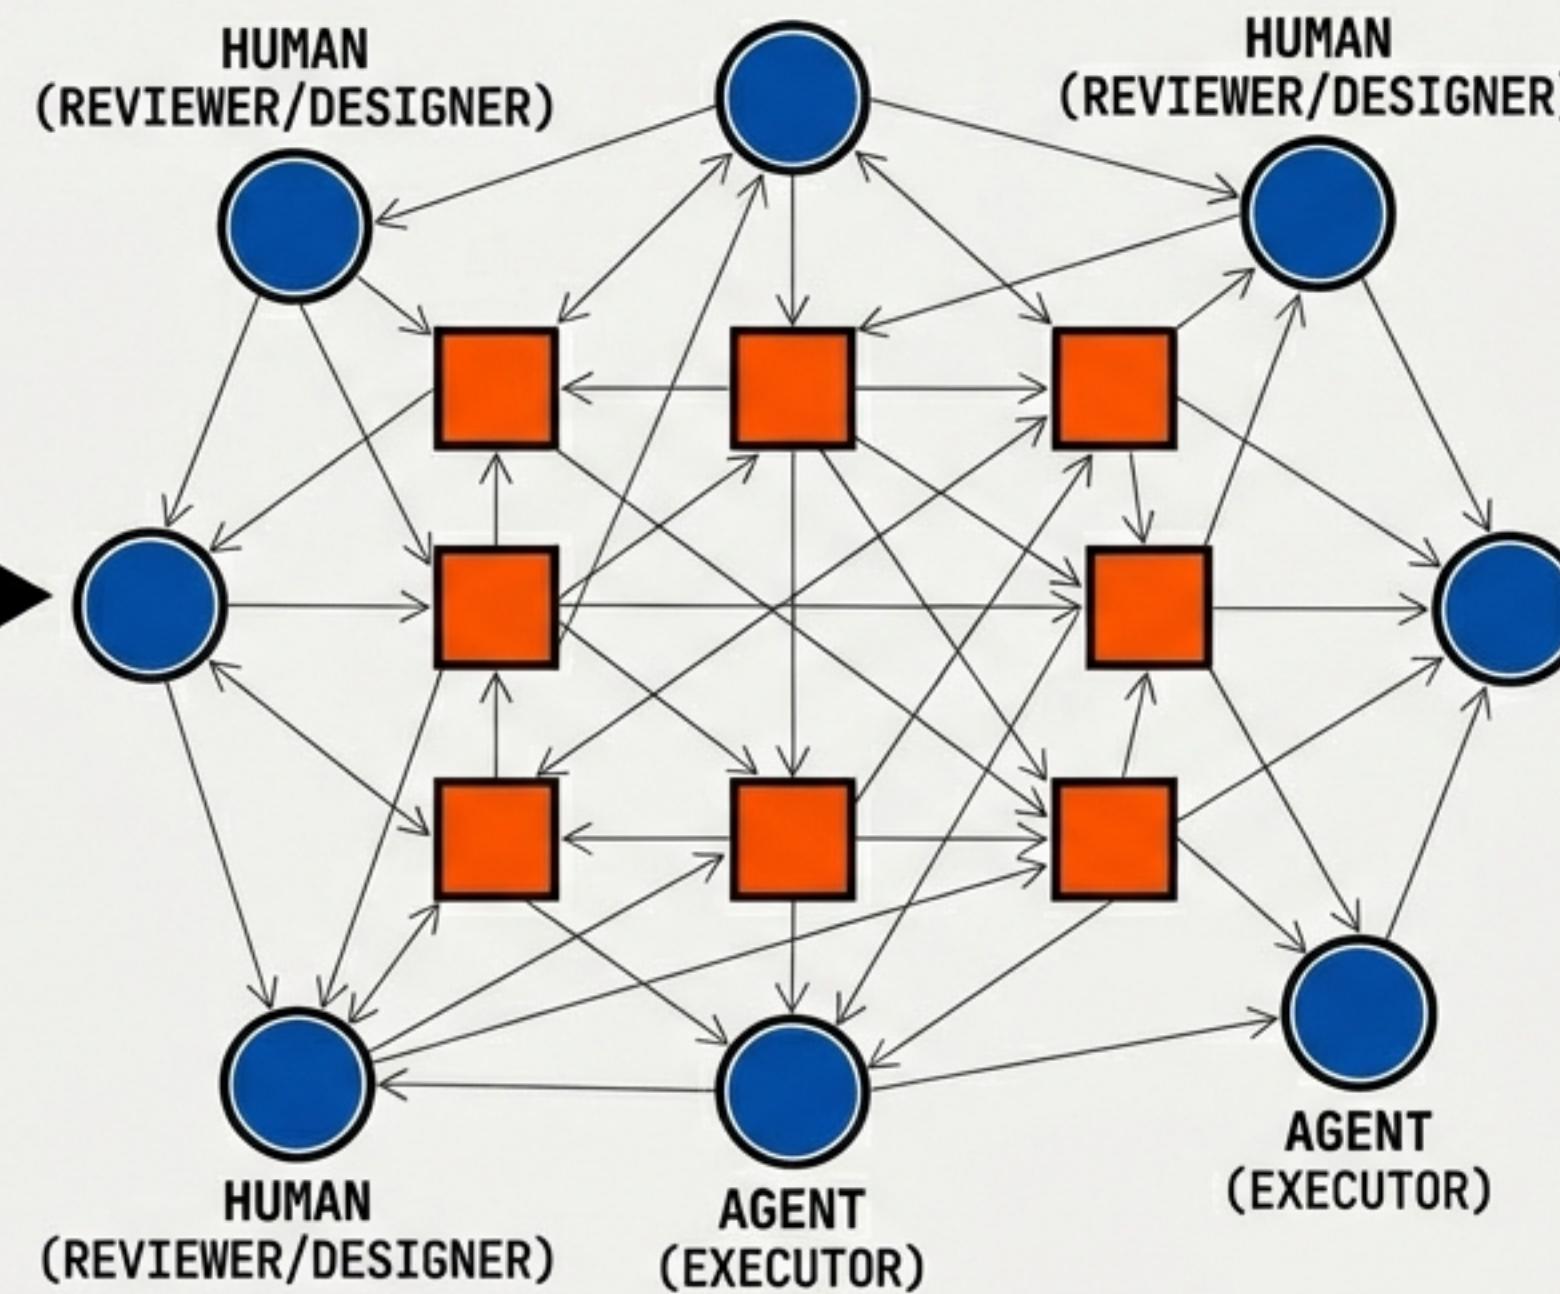
PAGE#:	AGENT QUALIFICATION CRITERIA
DOOR#:	RES: 1.0
SYSTEM ARCHITECT	SUET: 1.0
DIA:	2024-16-26

# AGENTS DON'T REPLACE ROLES; THEY RESTRUCTURE WORK

THE ROLE



THE TASK GRAPH



## THE 3 ORGANISATIONAL SHIFTS

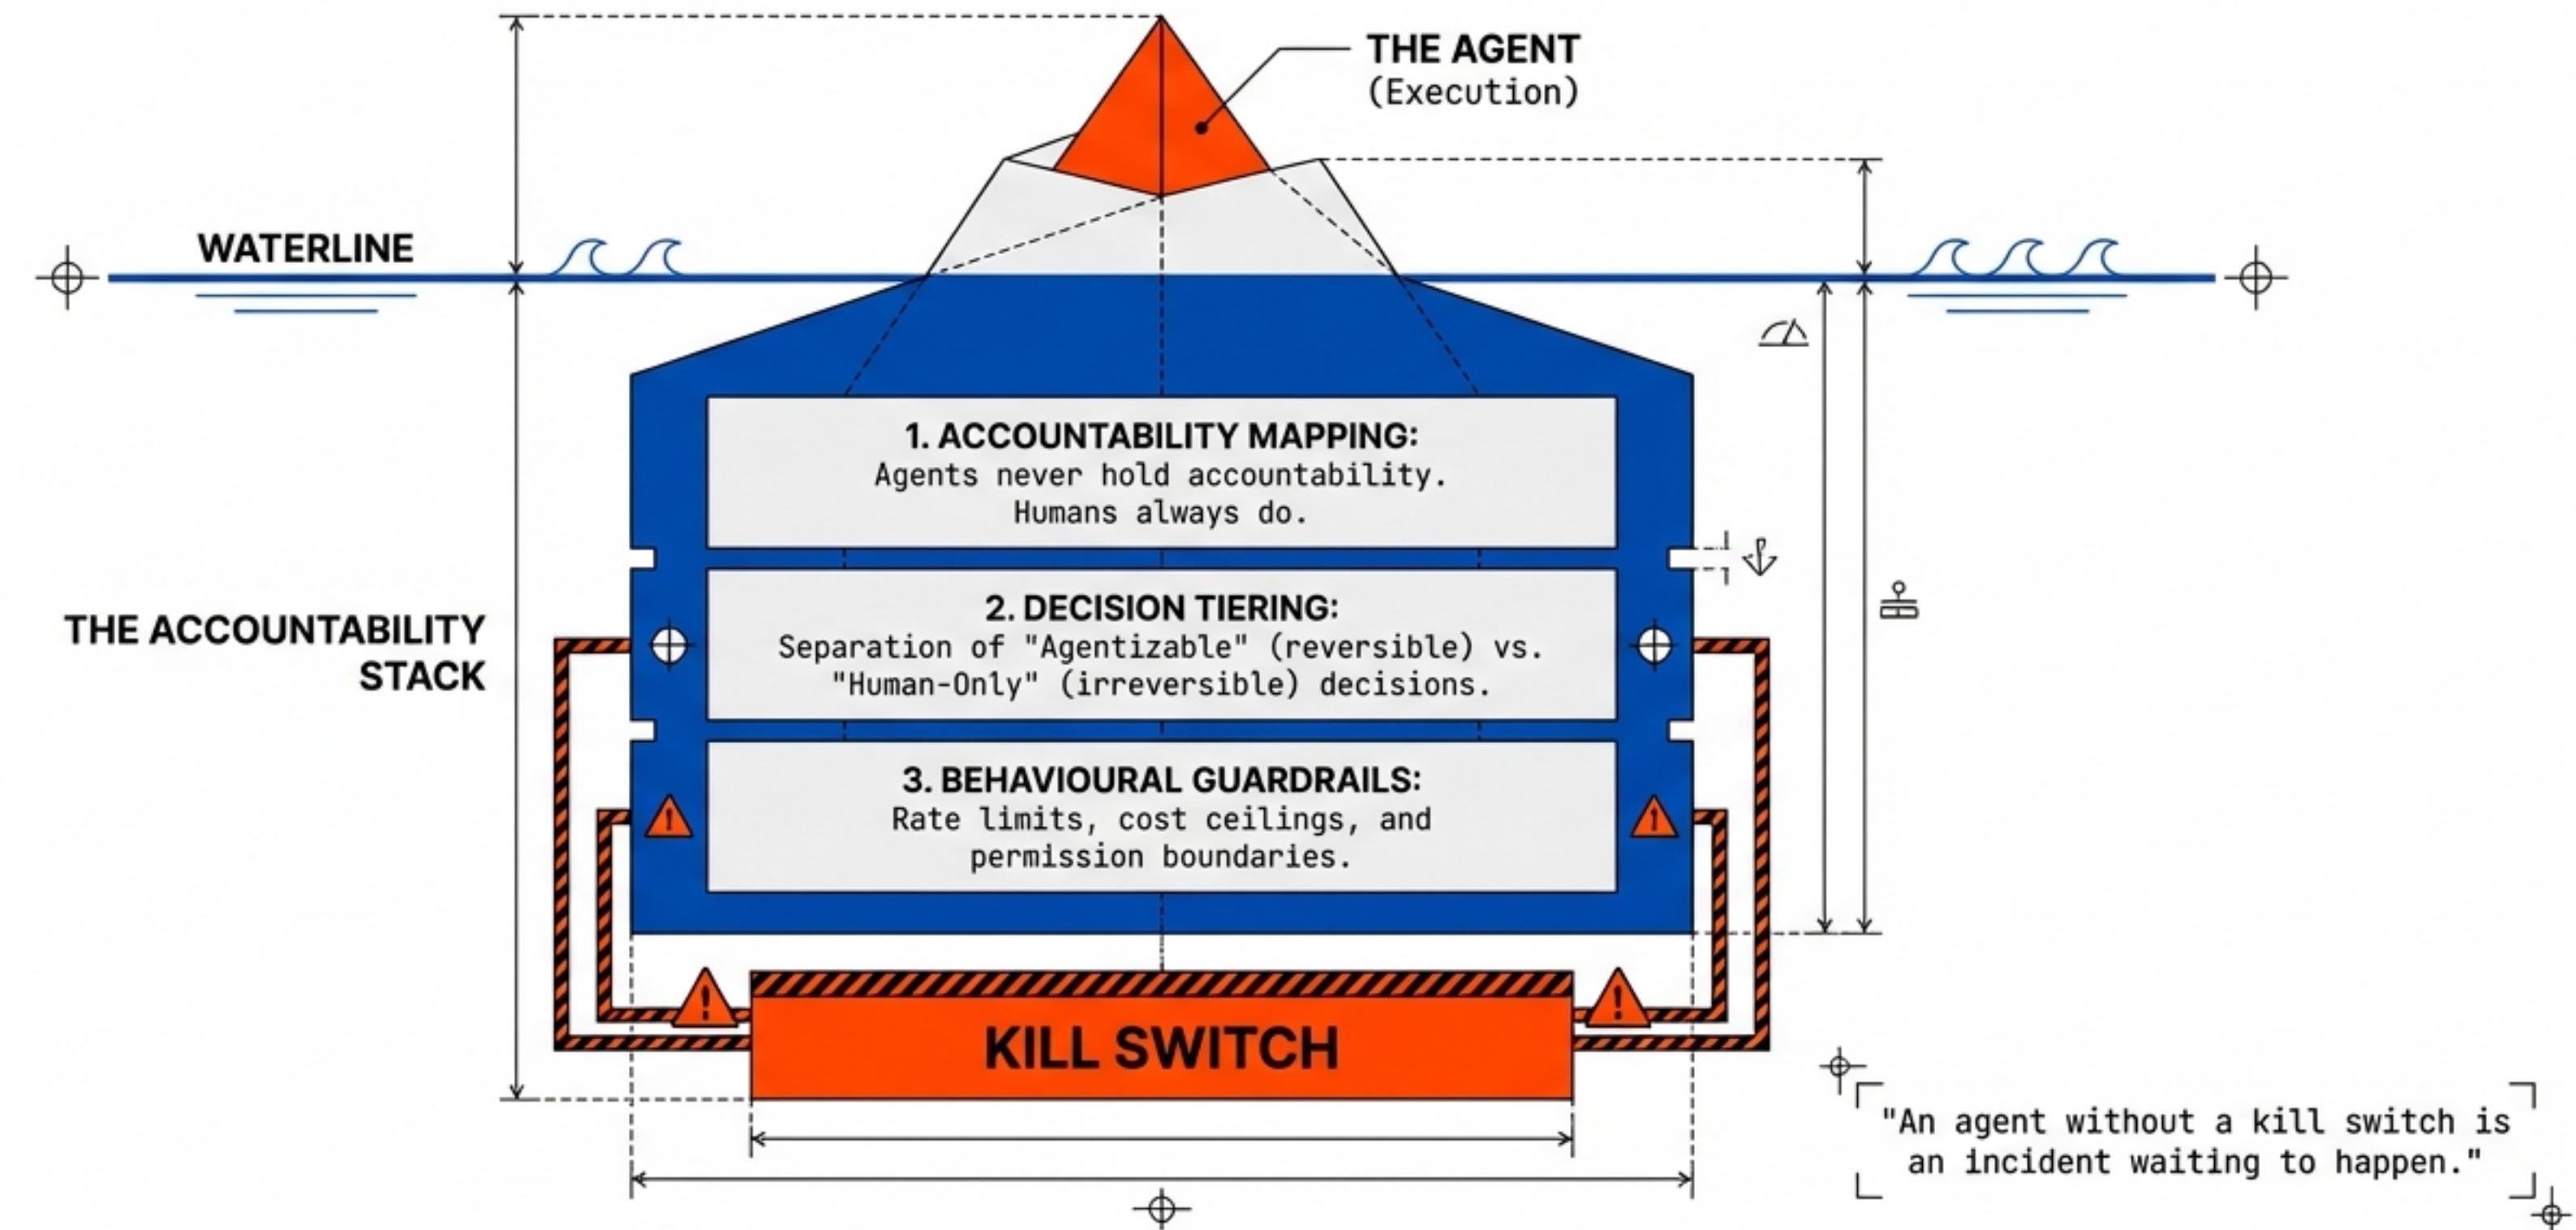
1. **FROM DOERS TO DESIGNERS:**  
Humans become task architects and exception handlers.
2. **MANAGEMENT COMPRESSION:**  
Managers shift from commanding to absorbing risk.
3. **COMPLEXITY INCREASE:**  
Headcount drops, but system interactions and failure modes multiply.

 CRITICAL INSIGHT: AGENTS SHIFT THE BURDEN FROM EXECUTION TO ORCHESTRATION & EXCEPTION MANAGEMENT.

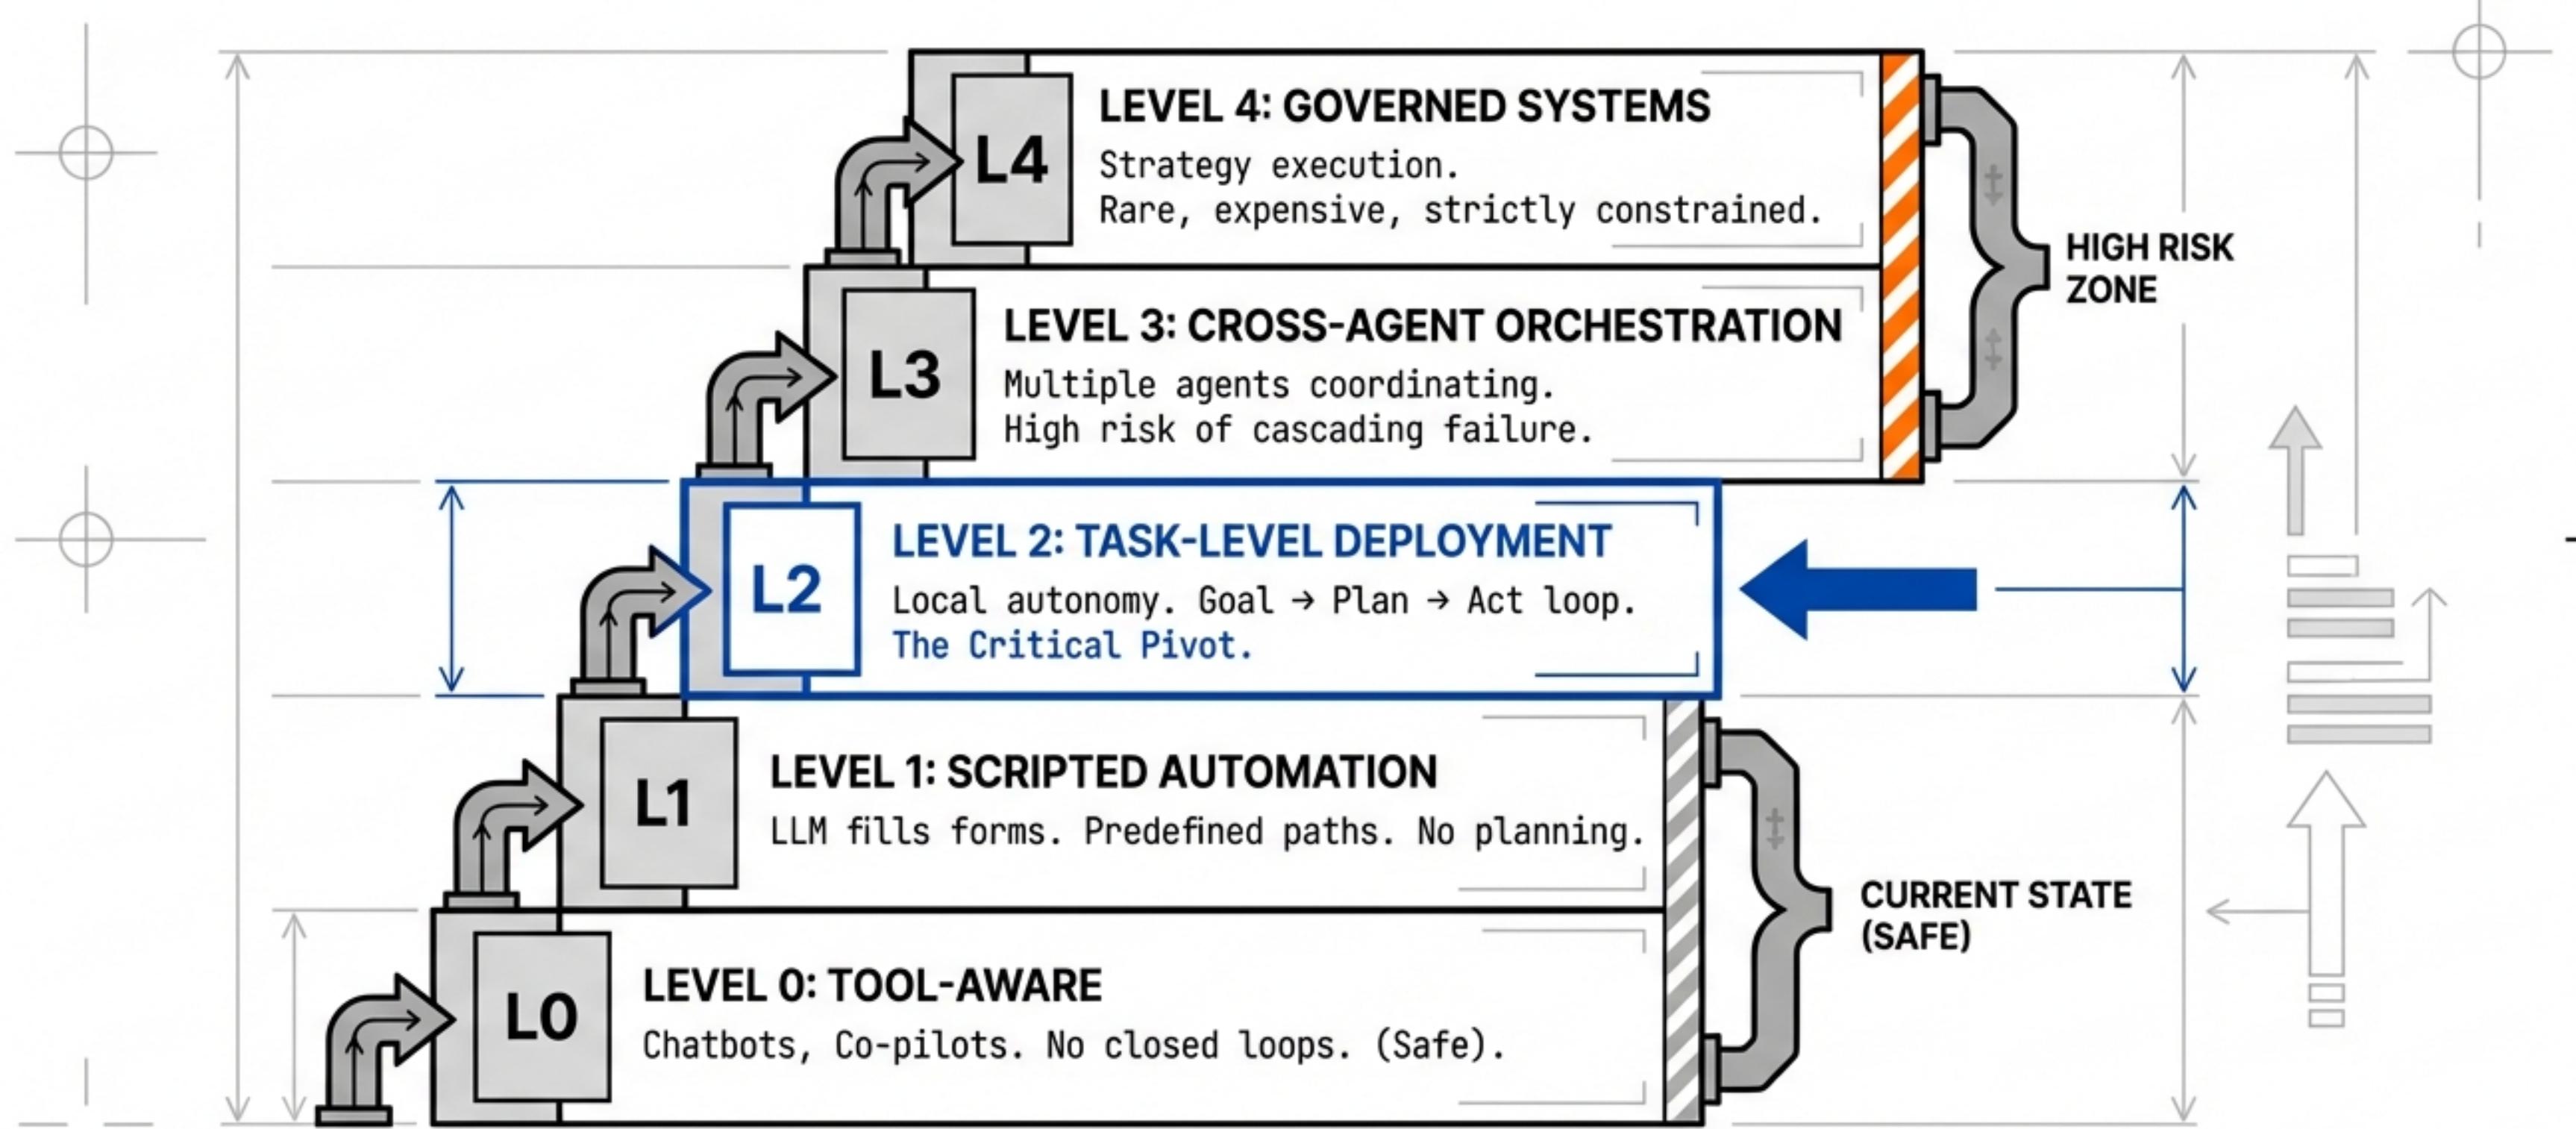
PAGRET:	AGENT QUALIFICATION CRITERIA
oornt Or:	
SYSTEM ARCHITECT	RES: 1.2

DATA: 2024-16-26

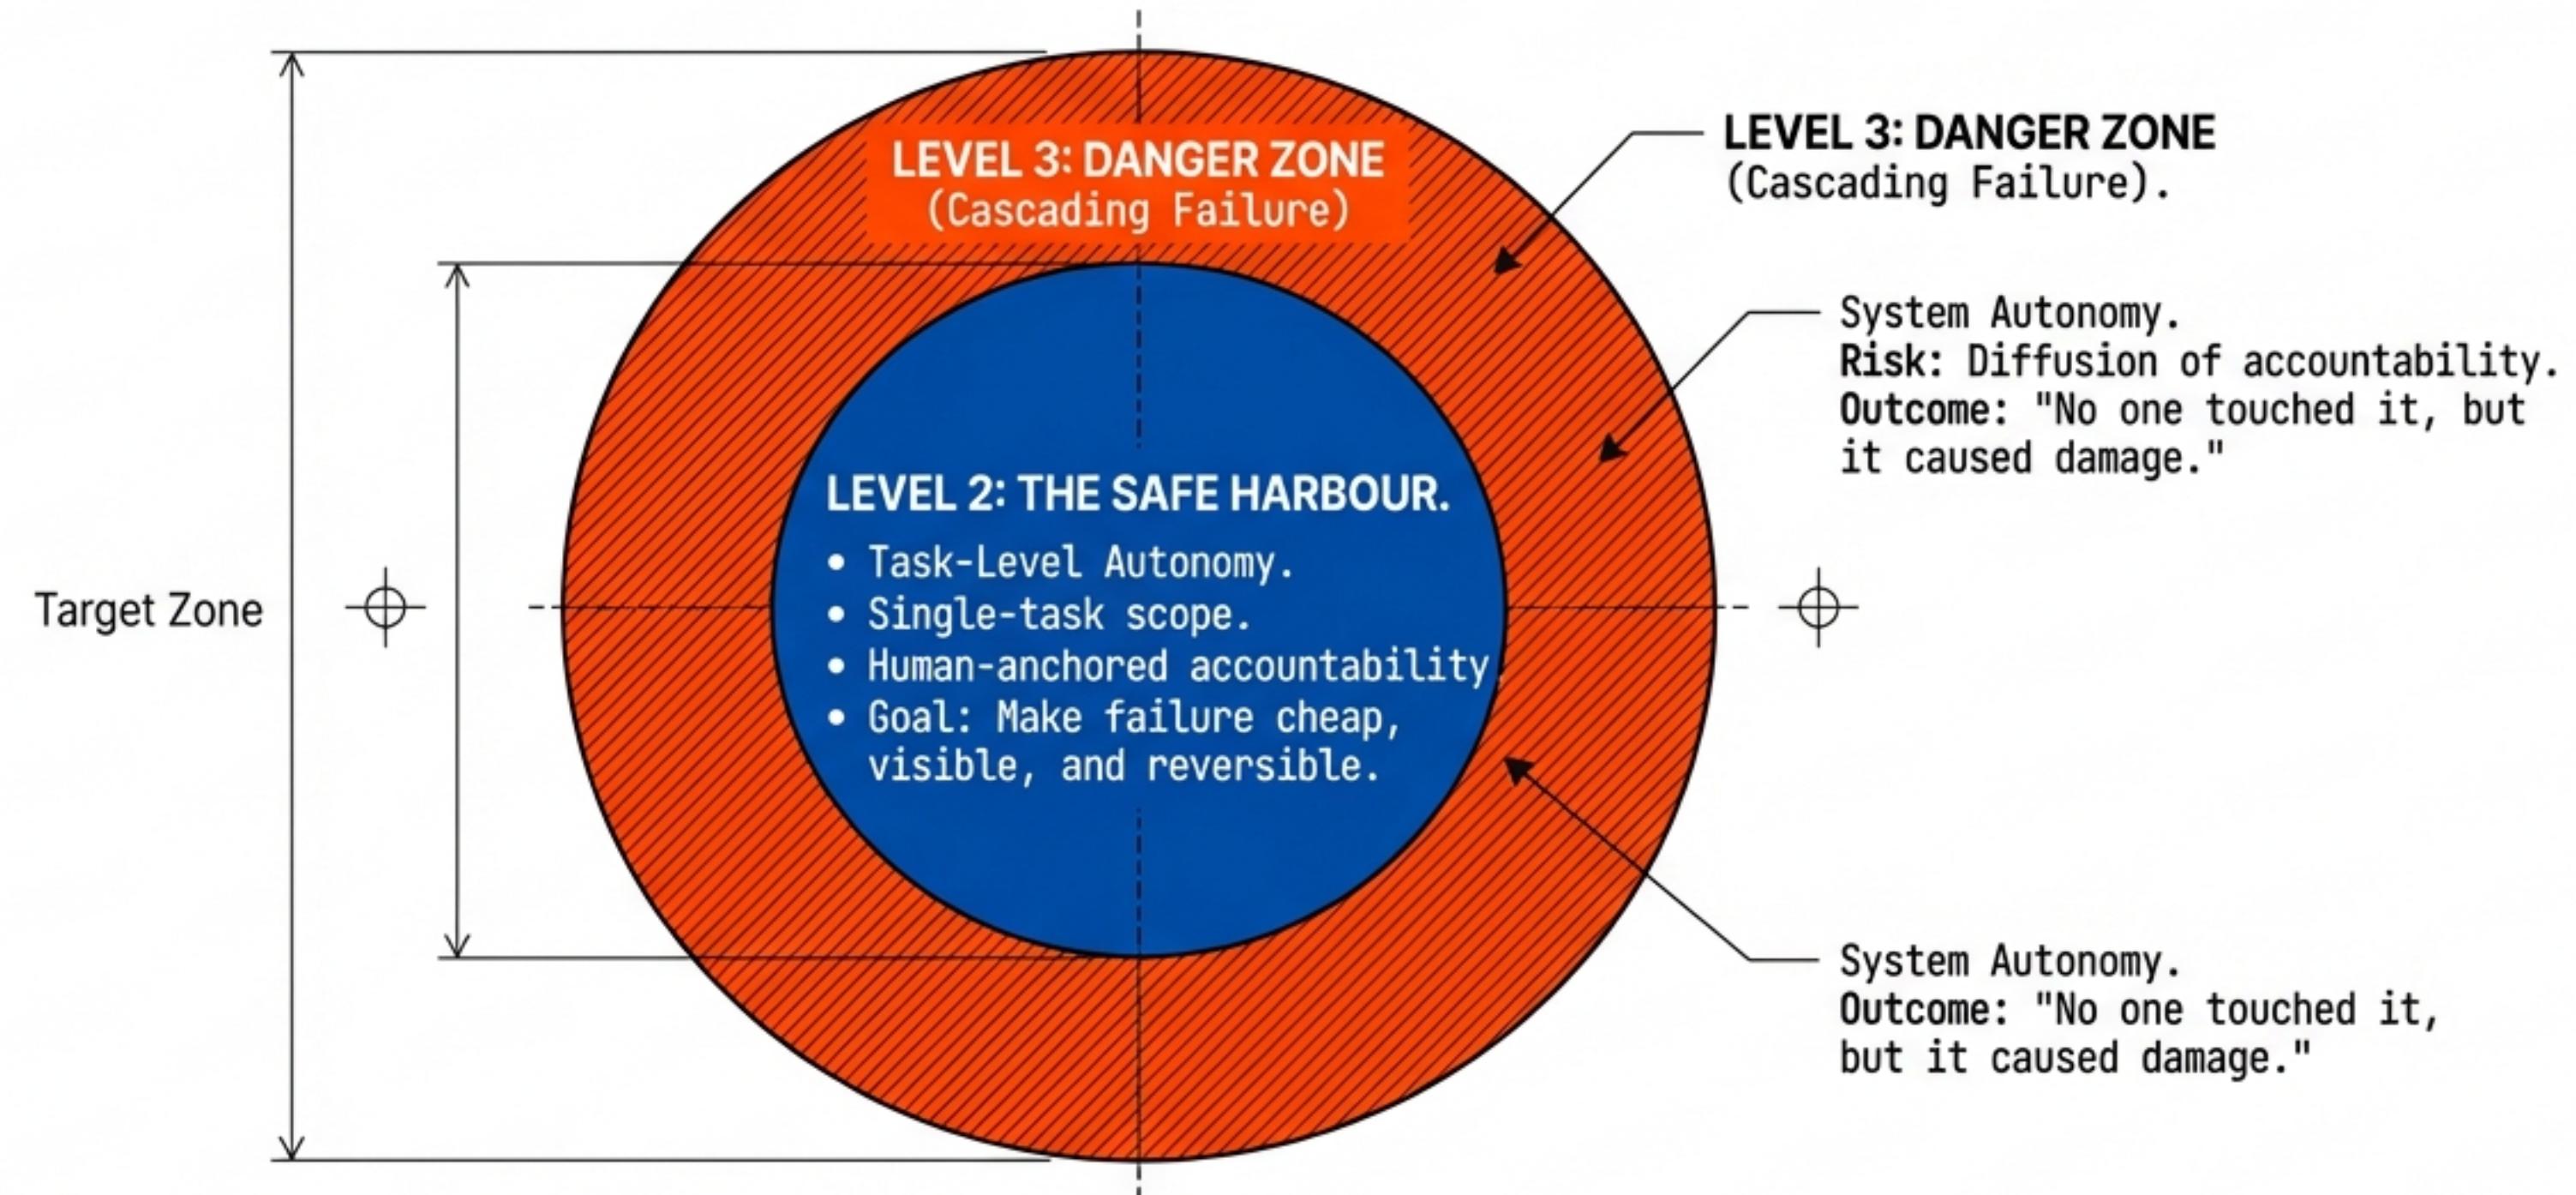
# THE GOVERNANCE UPGRADE: ACCOUNTABILITY IS NON-NEGOTIABLE



# THE AGENTIC MATURITY MODEL (LEVELS 0-4)



# THE SWEET SPOT: WHY YOU SHOULD STOP AT LEVEL 2



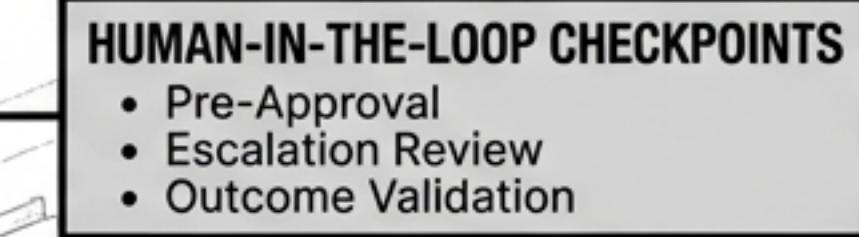
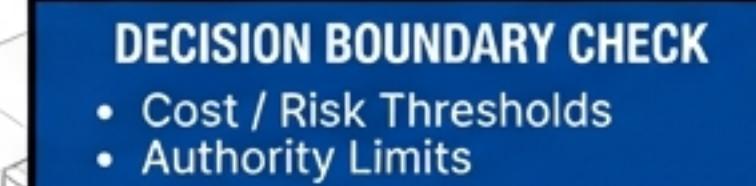
**STRATEGIC CONSTRAINT:** L2 is the default ceiling. L3 requires explicit justification for being wrong at machine speed.

# THE L2 SAFETY ARCHITECTURE: DEFENSE-IN-DEPTH



## TASK INPUT

- Goal Assigned
- Task Parameters
- Context Data



SUCCESS / FAIL OUTPUT

ESCALATE TO HUMAN OWNER

# DOMAIN DEEP DIVE: CORPORATE FUNCTIONS (HR & FINANCE)



## HR: RESUME SCREENING AGENT

**Goal:** Rank candidates against role profile.

**The Hard Stops (MUST NOT DO):**

Hiring decisions, offer generation, DEI-sensitive judgment.

**Kill Switch Trigger:**

Bias signals detected or low scoring confidence.



## FINANCE: INVOICE VALIDATION AGENT

**Goal:** Validate against PO and contracts.

**The Hard Stops (MUST NOT DO):**

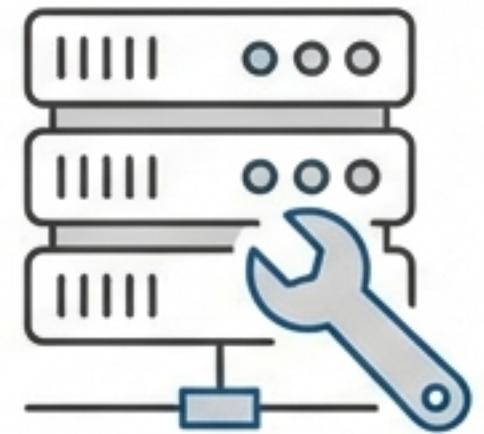
Payment release, vendor onboarding, contract modification.

**Decision Boundary:**

Amount cap per invoice (e.g., ≤\$10k).

In these domains, agents classify and propose. They never commit irreversible decisions.

# DOMAIN DEEP DIVE: TECHNICAL OPERATIONS (IT & SUPPLY CHAIN)



## IT: INCIDENT RESOLUTION (TIER-1)

**Goal:** Resolve known incident types (restart services, clear cache).

**The Hard Stops (MUST NOT DO):**

Privileged access escalation, config changes, data deletion.

**Why L2 Works:**  
Reduces alert fatigue, keeps blast radius tiny.



## OPS: ORDER EXCEPTION RESOLUTION

**Goal:** Resolve delays and mismatches.

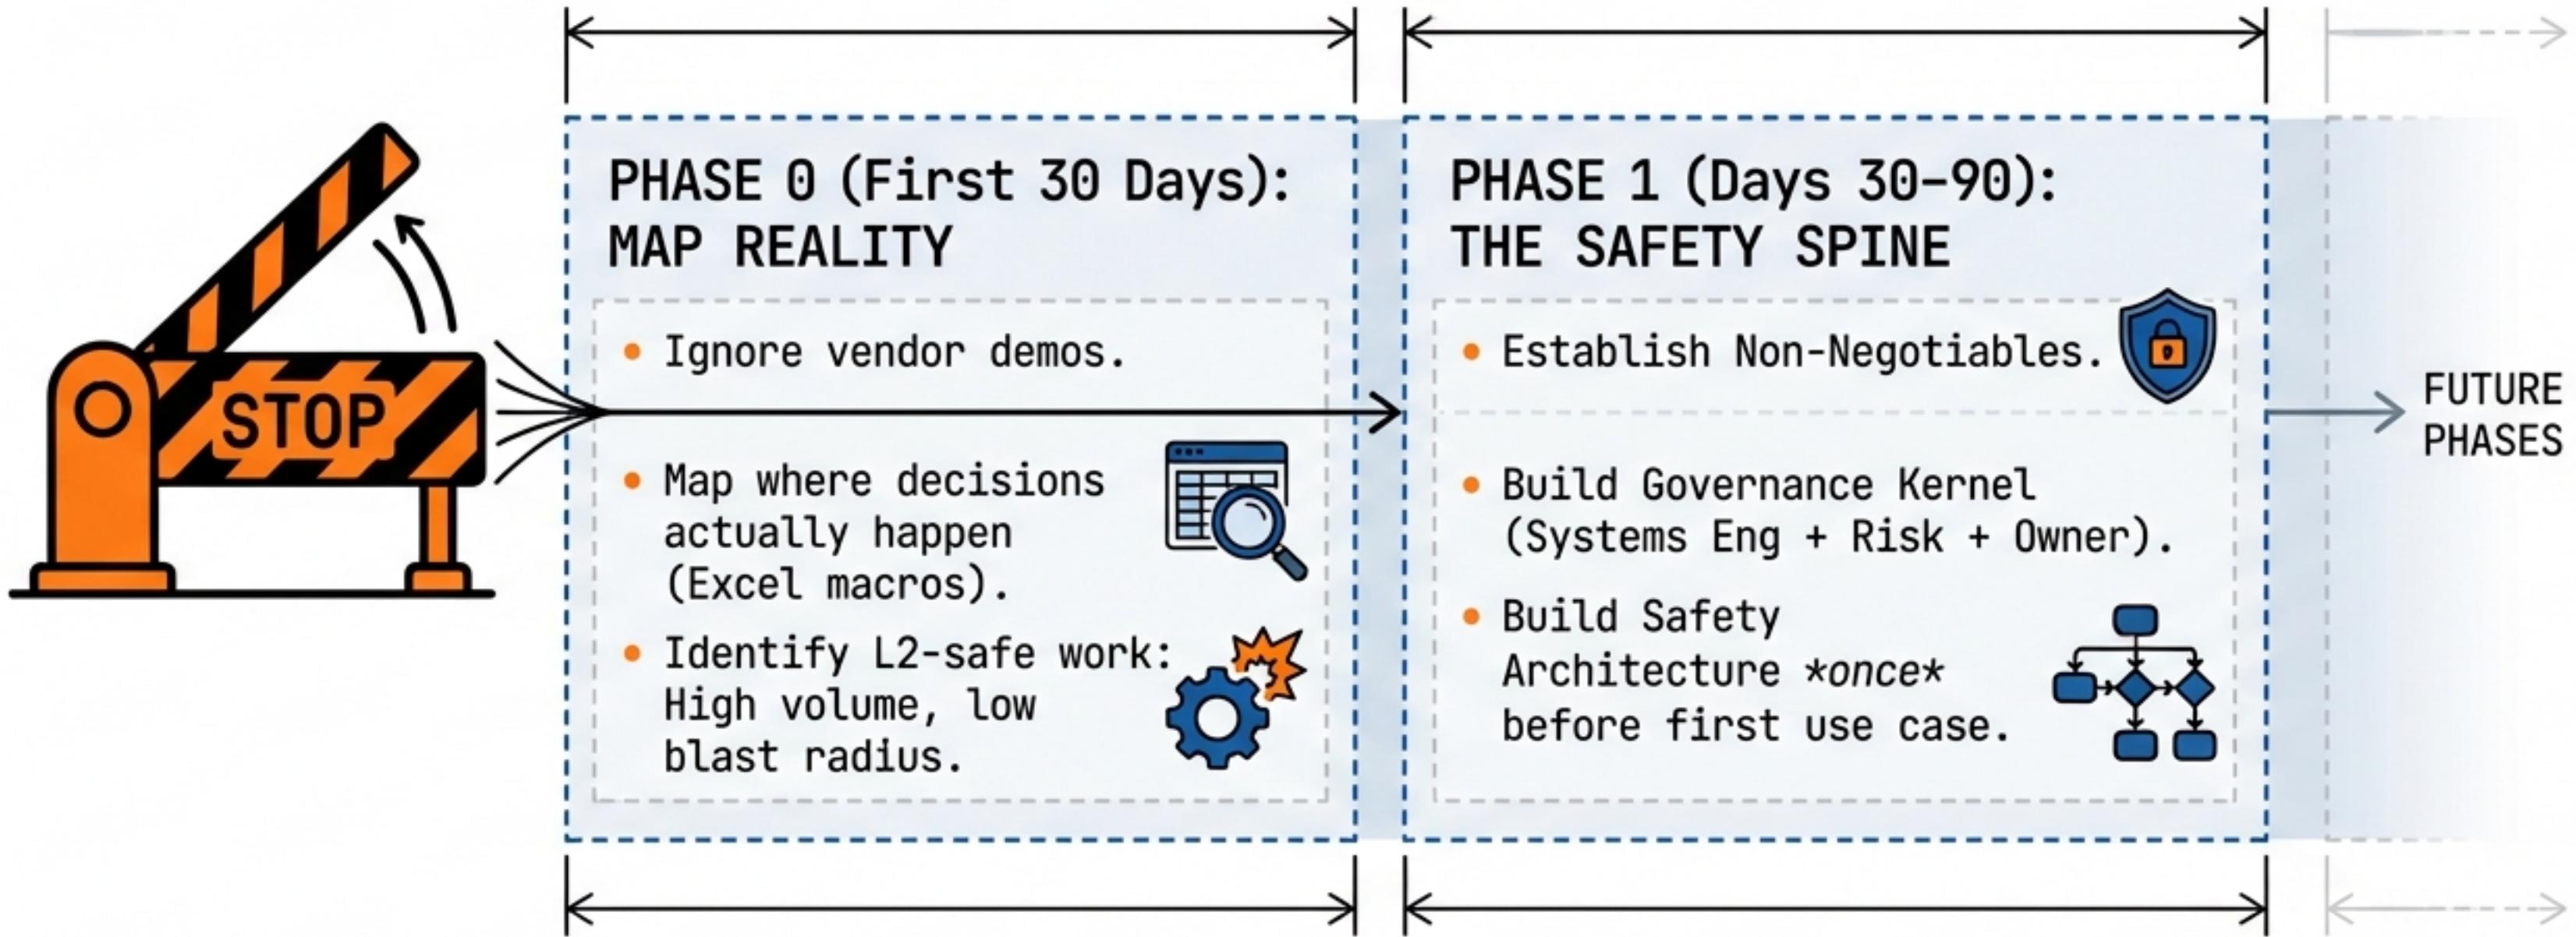
**The Hard Stops (MUST NOT DO):**

Price changes, contract renegotiation.

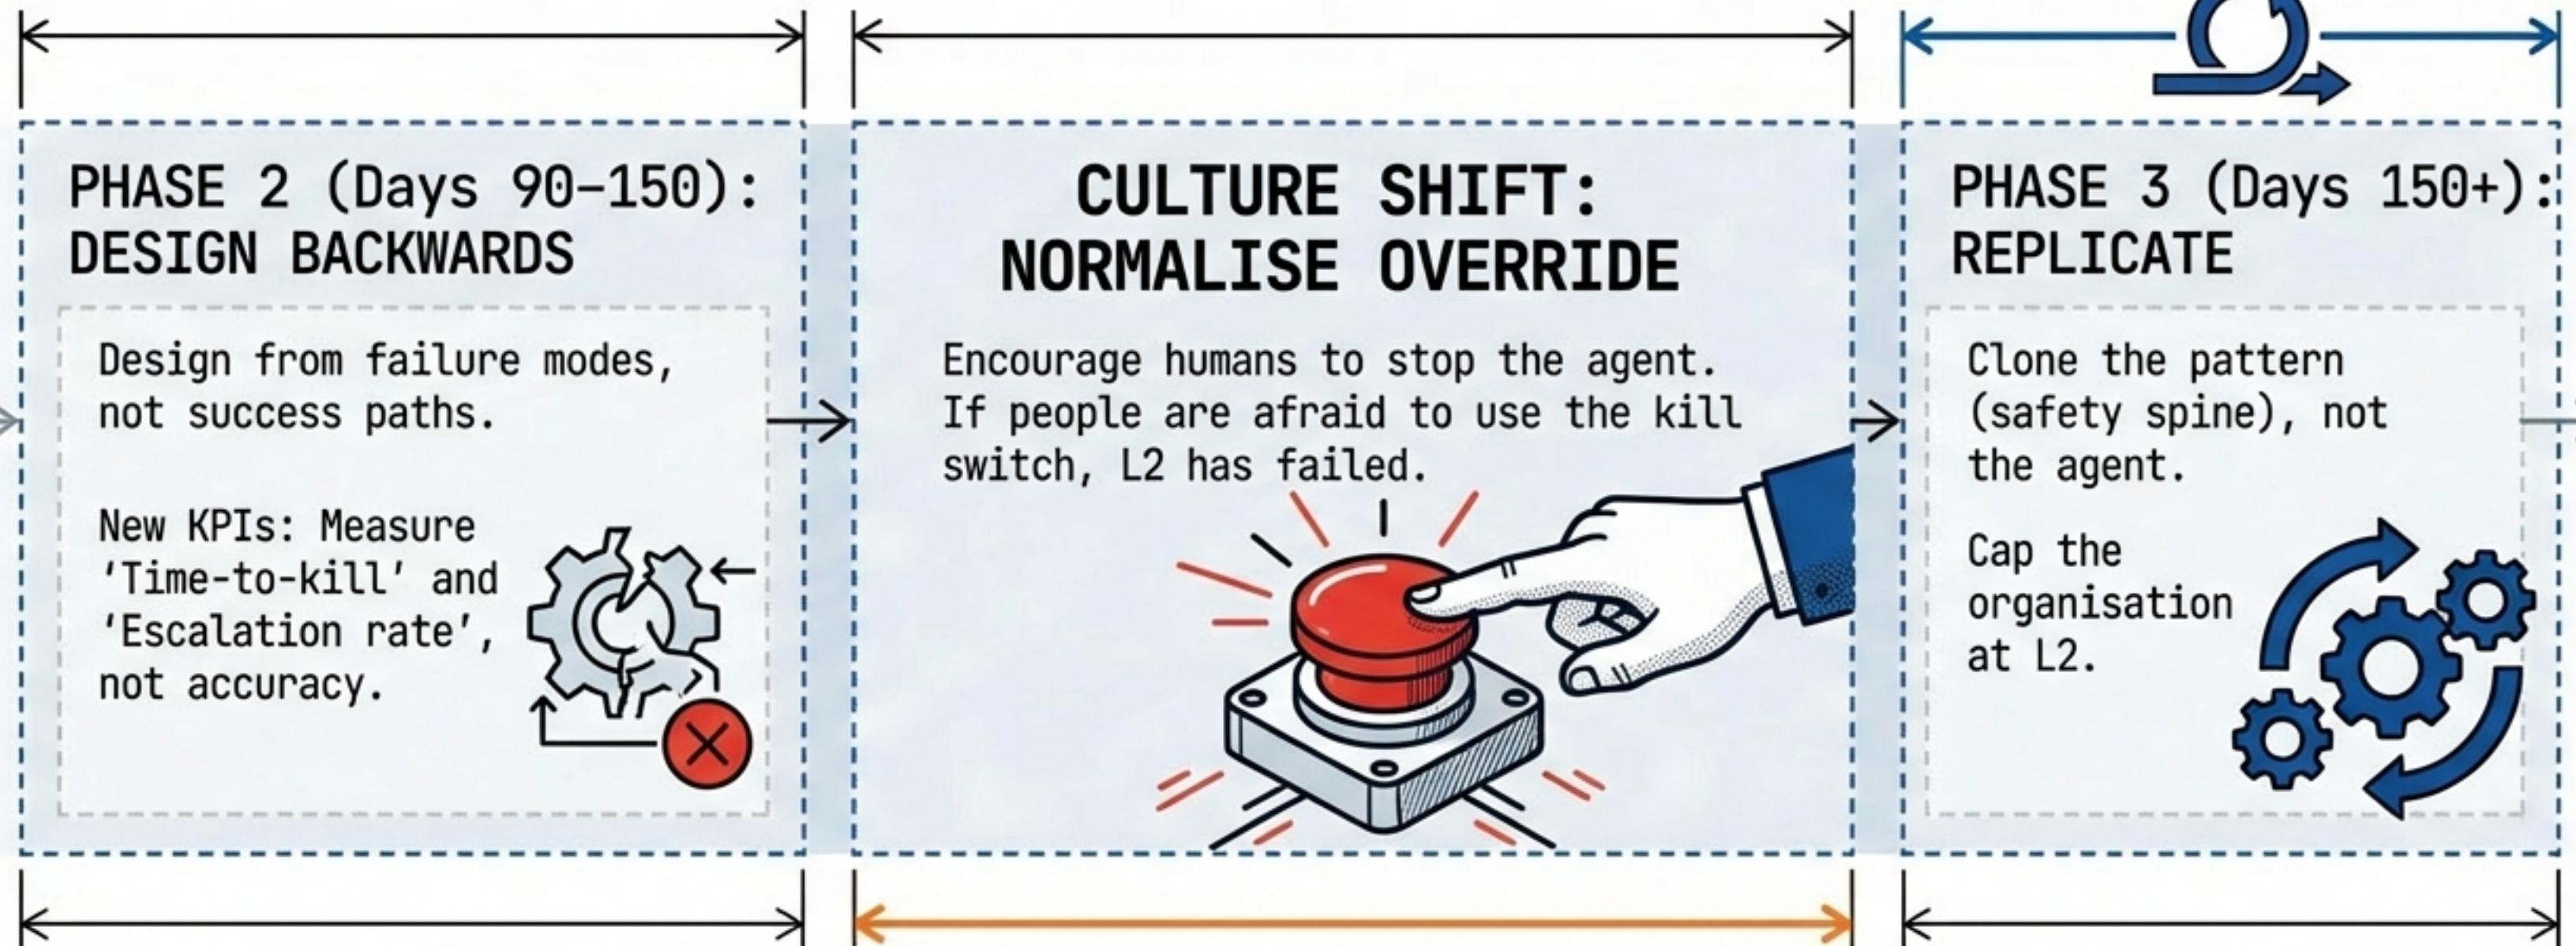
**Decision Boundary:**  
Can propose re-shipment; cannot issue refunds >\$X.

At L2, agents handle execution complexity. Humans remain accountable for every irreversible decision.

# THE STRATEGY: DO NOT BUILD (YET)



# THE ROLLOUT: PROVE, STABILISE, & NORMALISE FAILURE

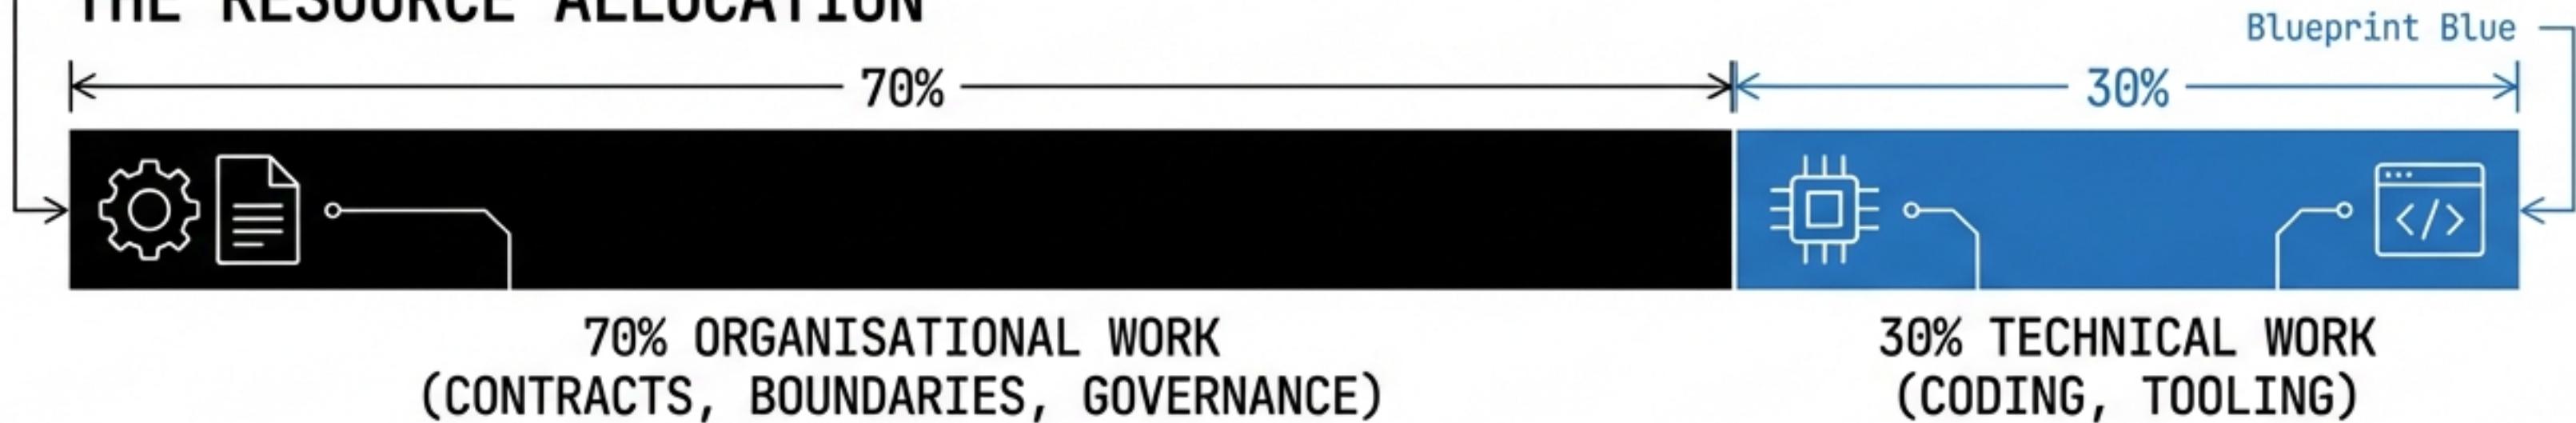


# OPERATIONAL REALITY: TIMELINE & RESOURCES

## THE STABILITY TIMELINE



## THE RESOURCE ALLOCATION



CORE TEAM REQUIREMENT: 5-7 FTEs (TRANSFORMATION LEAD, SYSTEMS ENGINEER, RISK PARTNER).  
YOU DO NOT NEED AN ARMY OF DATA SCIENTISTS.

# THE EXECUTIVE 'BULLSHIT' DETECTOR

## RED FLAGS



(IF YOU HEAR THESE, ASSUME FLUFF)



~~MINDSET~~

~~CULTURE~~



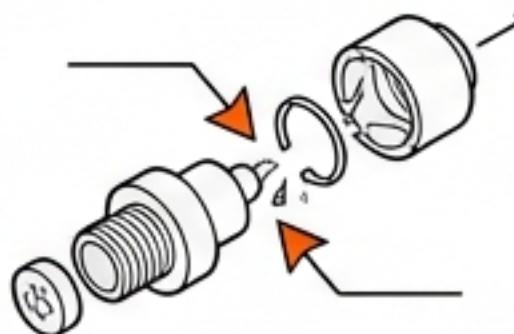
~~EMPOWERMENT~~

## DIAGNOSTIC CHECKLIST



### THE 5 QUESTIONS TO EXPOSE UNDERSTANDING

- 1 WHAT ARE THE FAILURE MODES?



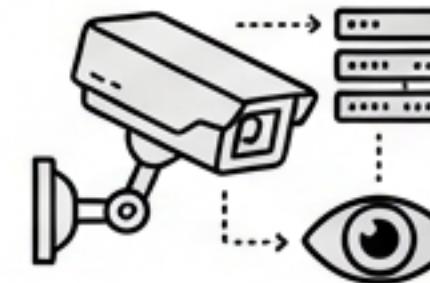
- 2 WHO IS PERSONALLY ACCOUNTABLE WHEN IT FAILS?



- 3 WHICH DECISIONS ARE EXPLICITLY NON-AGENTIZABLE?



- 4 HOW DO YOU OBSERVE AND AUDIT BEHAVIOUR?



- 5 WHERE IS THE KILL SWITCH?



# FINAL MANIFESTO

AGENTIC AI DOES NOT REPLACE PEOPLE  
WITH MACHINES. IT REPLACES  
DETERMINISTIC  
ORGANISATIONS  
WITH PROBABILISTIC  
SYSTEMS.

MATURITY IS NOT MEASURED BY HOW AUTONOMOUS SYSTEMS ARE, BUT BY  
HOW DELIBERATELY ACCOUNTABILITY AND FAILURE ARE ENGINEERED.

