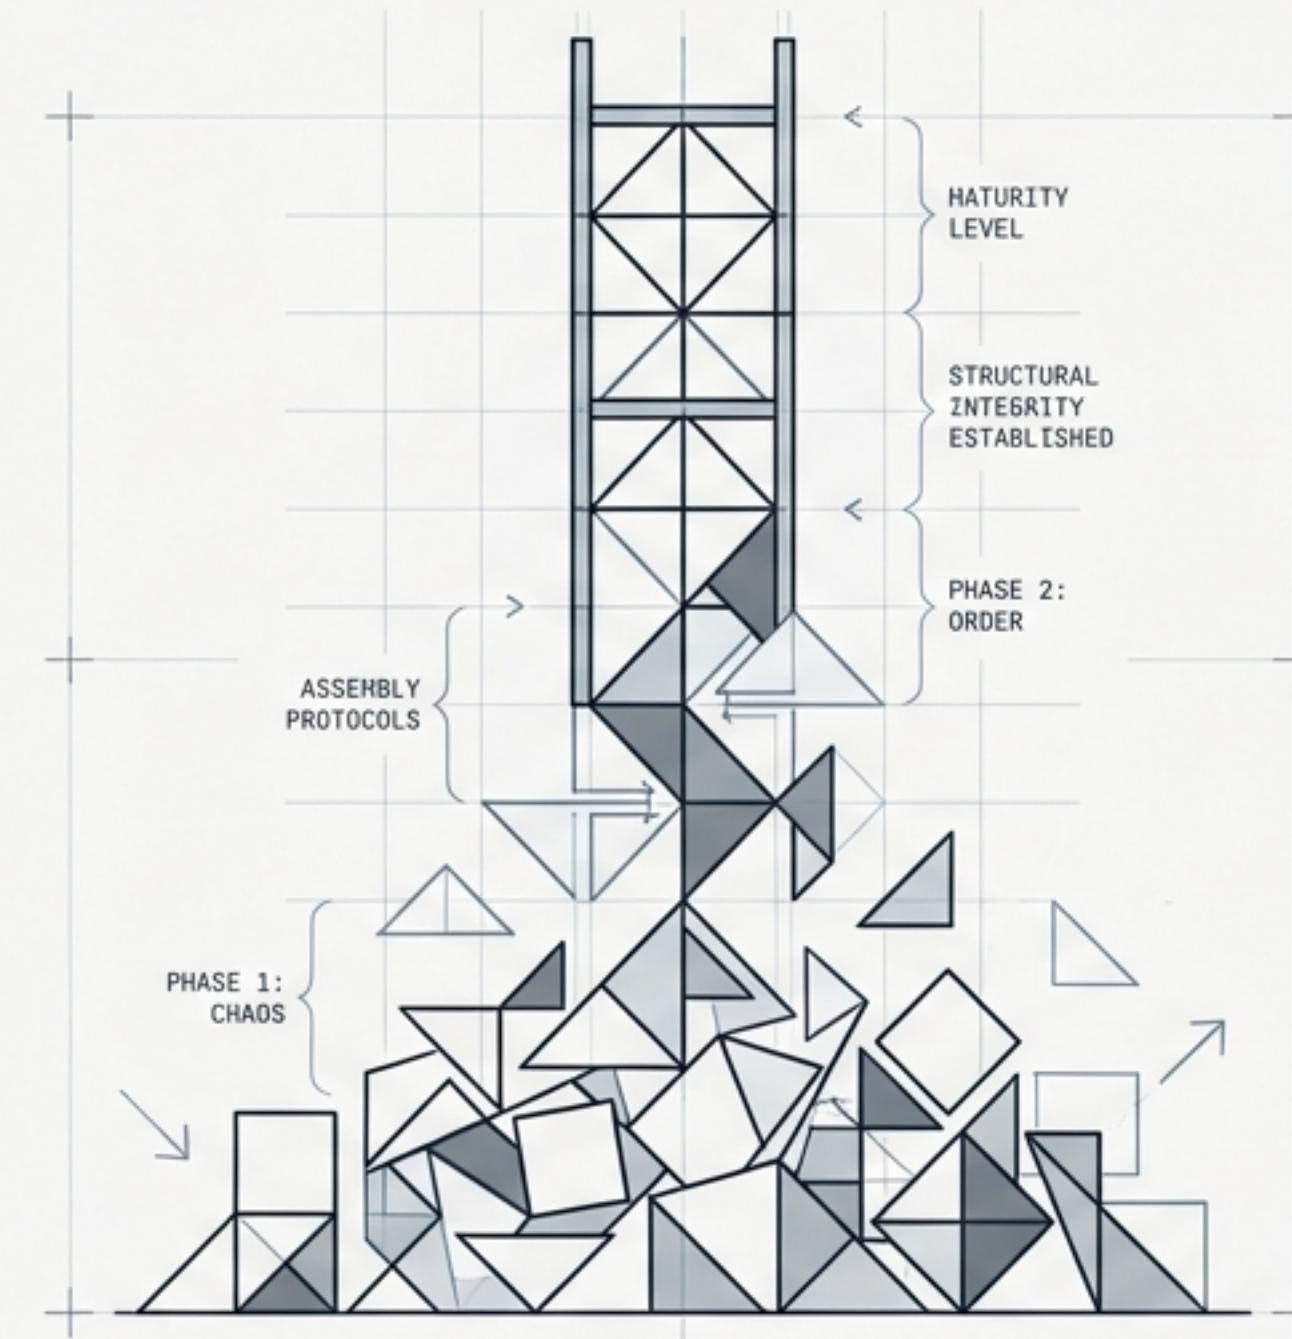


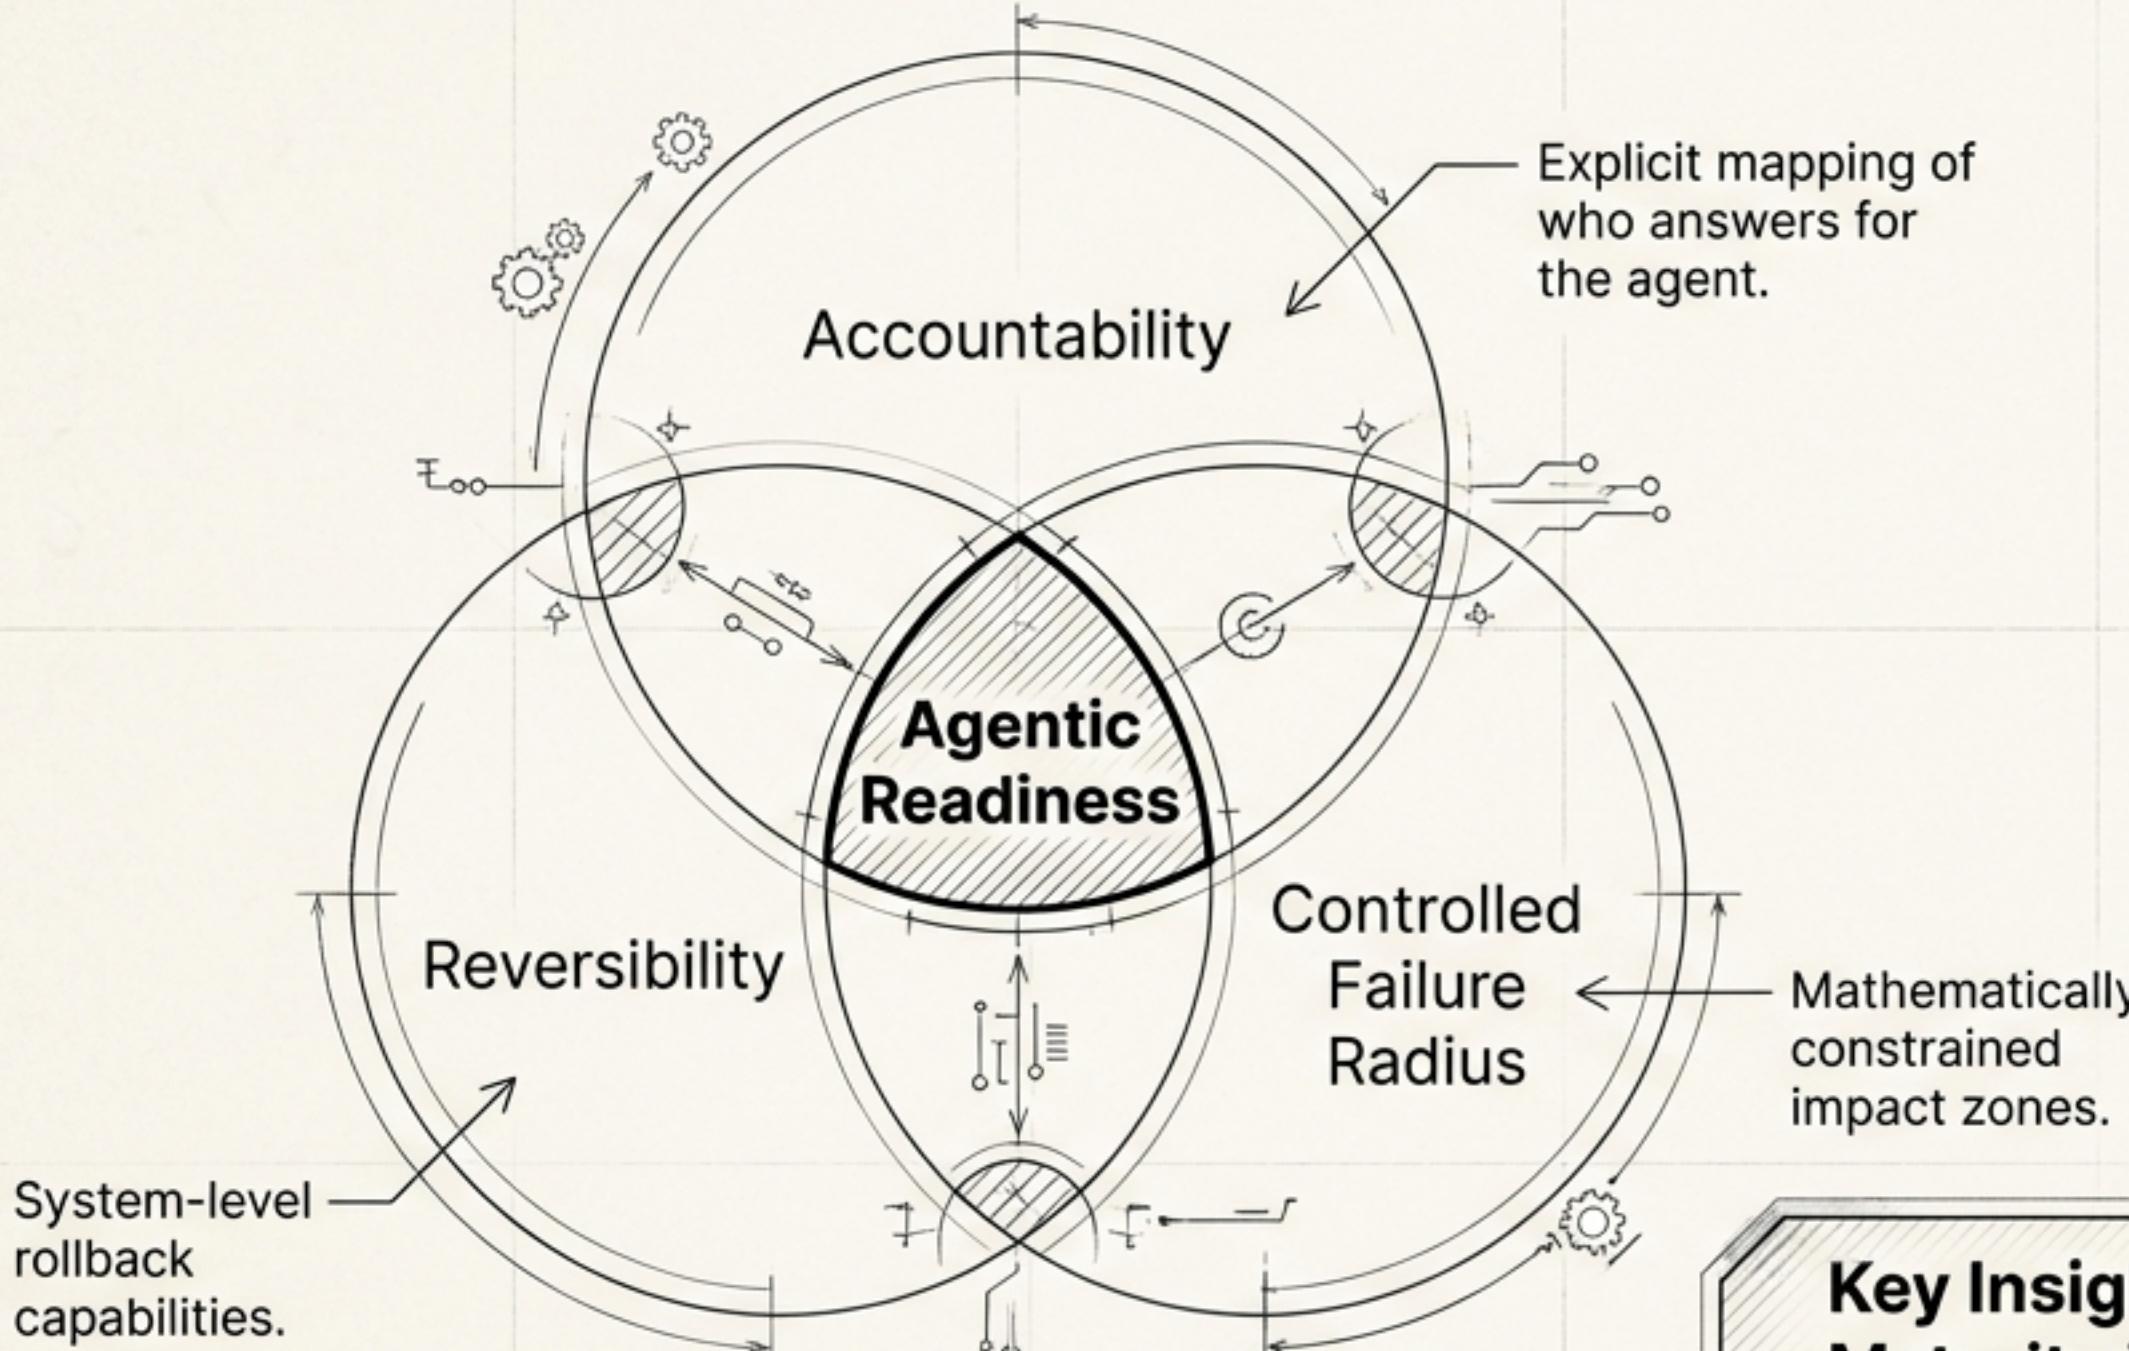
The Agentic Organisation Maturity Model

A Practical Guide to Moving Beyond the Hype of AI Autonomy



The Foundational Truth: Redefining Success

Stop measuring sophistication; start measuring safety.



Core Definition

An **agentic organisation** is not defined by its use of AI tools, but by its deliberate **design to delegate bounded execution to probabilistic systems**.

Key Insight:
Maturity is not about autonomy level; it's about how explicitly failure is designed.

LEVEL 0 — The Tool-Aware Organisation (The Baseline)



THE EXECUTIVE TELL

“We encourage our teams to experiment with AI tools.”

THE REALITY CHECK

Status: Where 80% of companies reside.



Technology

LLMs function as advanced assistants (chatbots, writing aids, code copilots). Interactions are shallow.



Governance

Basic data leakage policies. AI usage is treated like standard SaaS.

The Danger: Overestimating maturity by confusing ‘usage’ with ‘capability’.

Action Required

Action: Acknowledge this as foundation-building. Prohibit language that inflates capability.

LEVEL 1 — Scripted Automation with LLM Augmentation

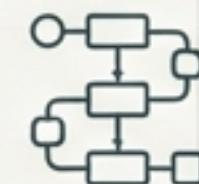


THE EXECUTIVE TELL

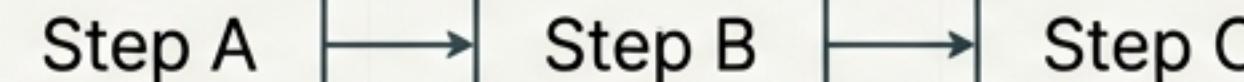
AI accelerates our existing processes.

THE REALITY CHECK

Status: The plateau of most “transformations”.



Technology



Deterministic workflows. The LLM is a component, not an actor.



Governance

Humans remain primary operators. Clear human accountability exists.

The Danger: Thinking this brittle architecture will scale to complex decisions.

Action Required

Action: Document process boundaries and identify decision points before attempting to scale.

LEVEL 2 — Task-Level Agent Deployment (The Danger Zone)

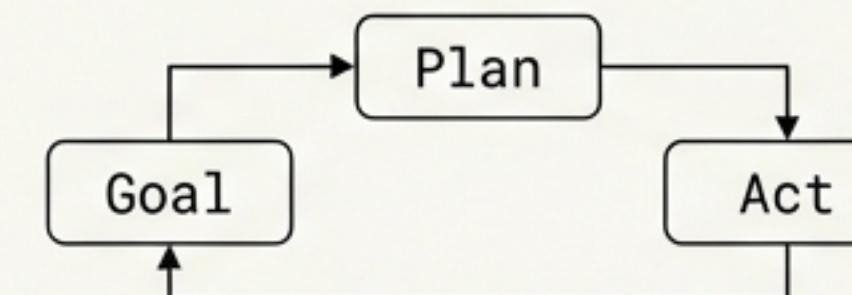
THE EXECUTIVE TELL

Agents handle routine decisions, humans handle judgment.

THE REALITY CHECK

Status: High-Risk Pilot Zone.

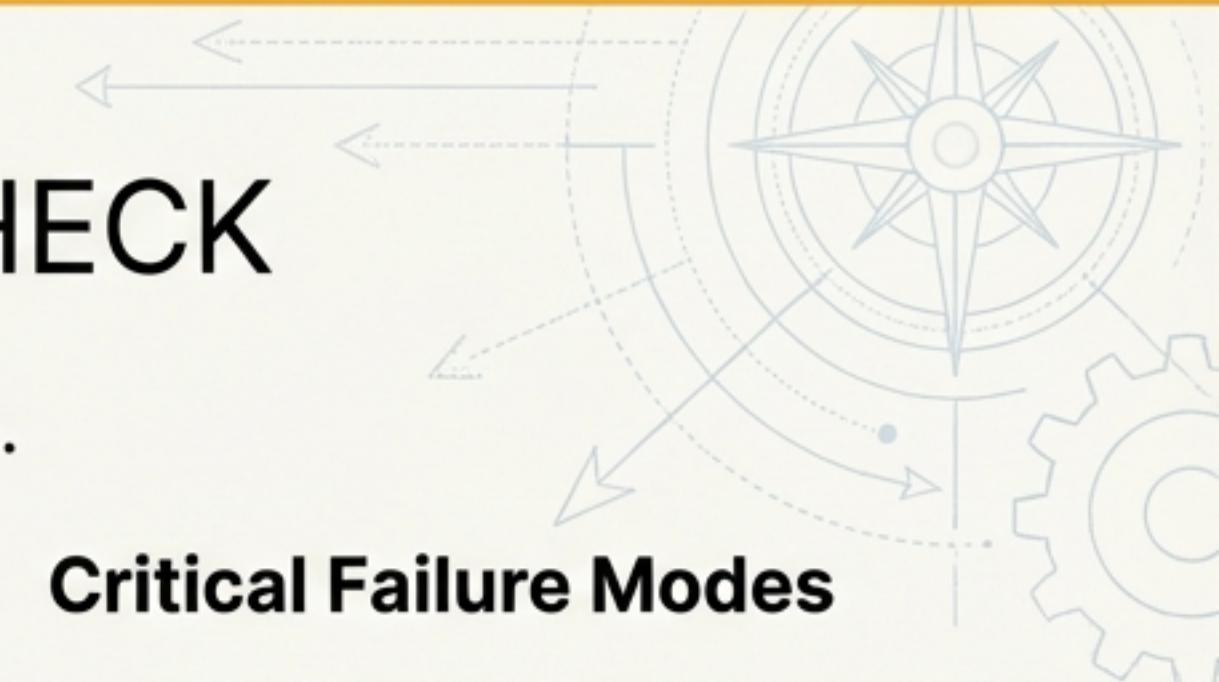
Technology



Real agents operating in narrow domains. Humans shift from operators to exception handlers.

Action Required

Action: Implement “Negative Space” rules (what the agent CANNOT do) and mandatory pre-mortems.



Critical Failure Modes

- **Silent Scope Creep:** “Let the agent handle just a bit more”.
- **Accountability Erosion:** Who owns the agent’s output?

LEVEL 3 — Cross-Agent Orchestration (The Complexity Cliff)



THE EXECUTIVE TELL

“Our operating model is becoming AI-orchestrated.”

— CEO / CTO

THE REALITY CHECK

Status: Only for engineering excellence.

Technology

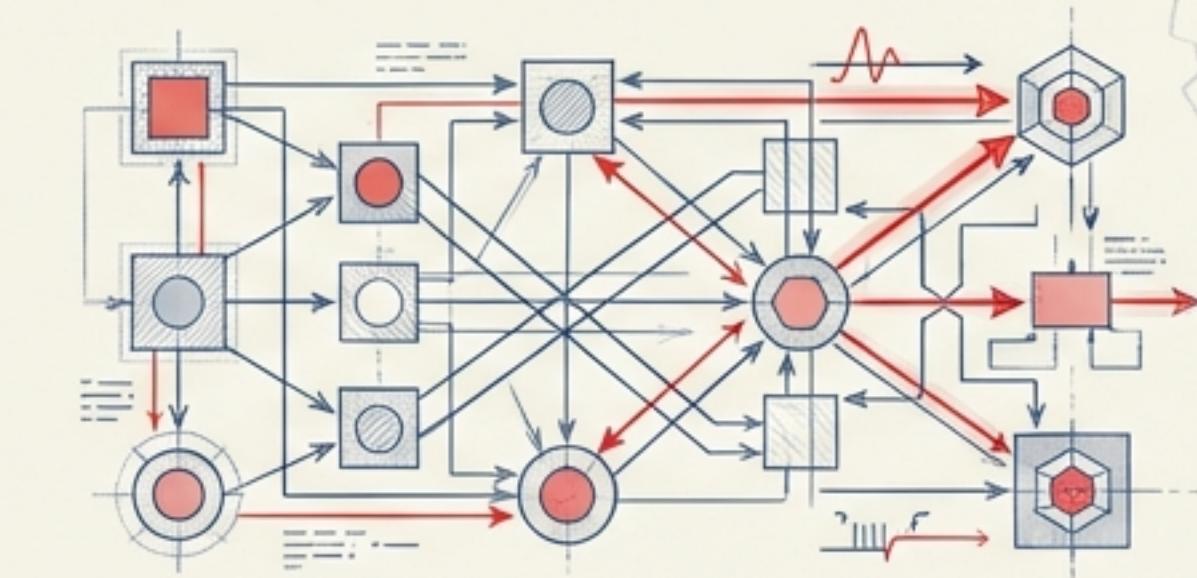
Emergent behaviours appear.
Multiple agents coordinate across functions.

The Danger:

Cascading Failures. One agent's error propagates through the system at speed.

Action Required

Action: Require “Circuit Breakers” (auto-isolation) and ‘Causation Tracing’ before entering this level.



LEVEL 4 — Governed Autonomous Decision Systems



THE EXECUTIVE TELL

We design
which
uncertainty
we accept.

THE REALITY CHECK

Status: Mathematically bounded autonomy. Rare.



Technology

Humans intervene by architectural design, not exception. Failure radius is mathematically constrained.

The Constraints

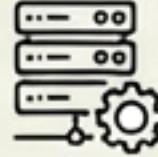
- High Cost
- Extreme Talent Density
- Regulatory Scrutiny

Action Required

Action: Reserve only for existential competitive advantages (e.g., high-frequency trading).

The Strategic Pivot: Why Higher Isn't Always Better

Concept: The Maturity Ceiling

Strategic Constraints		
 Function	 Maturity Cap	 Constraint
 Finance & Legal	 LEVEL 2	Mandatory human review for external actions.
 Marketing	 LEVEL 3	A/B testing allowed with strict spend caps.
 IT Operations	 LEVEL 3	Incident response allowed; commander's intent preserved.

Strategic constraint is a feature, not a bug.

Diagnostic Honesty: The ‘Bullshit Filter’



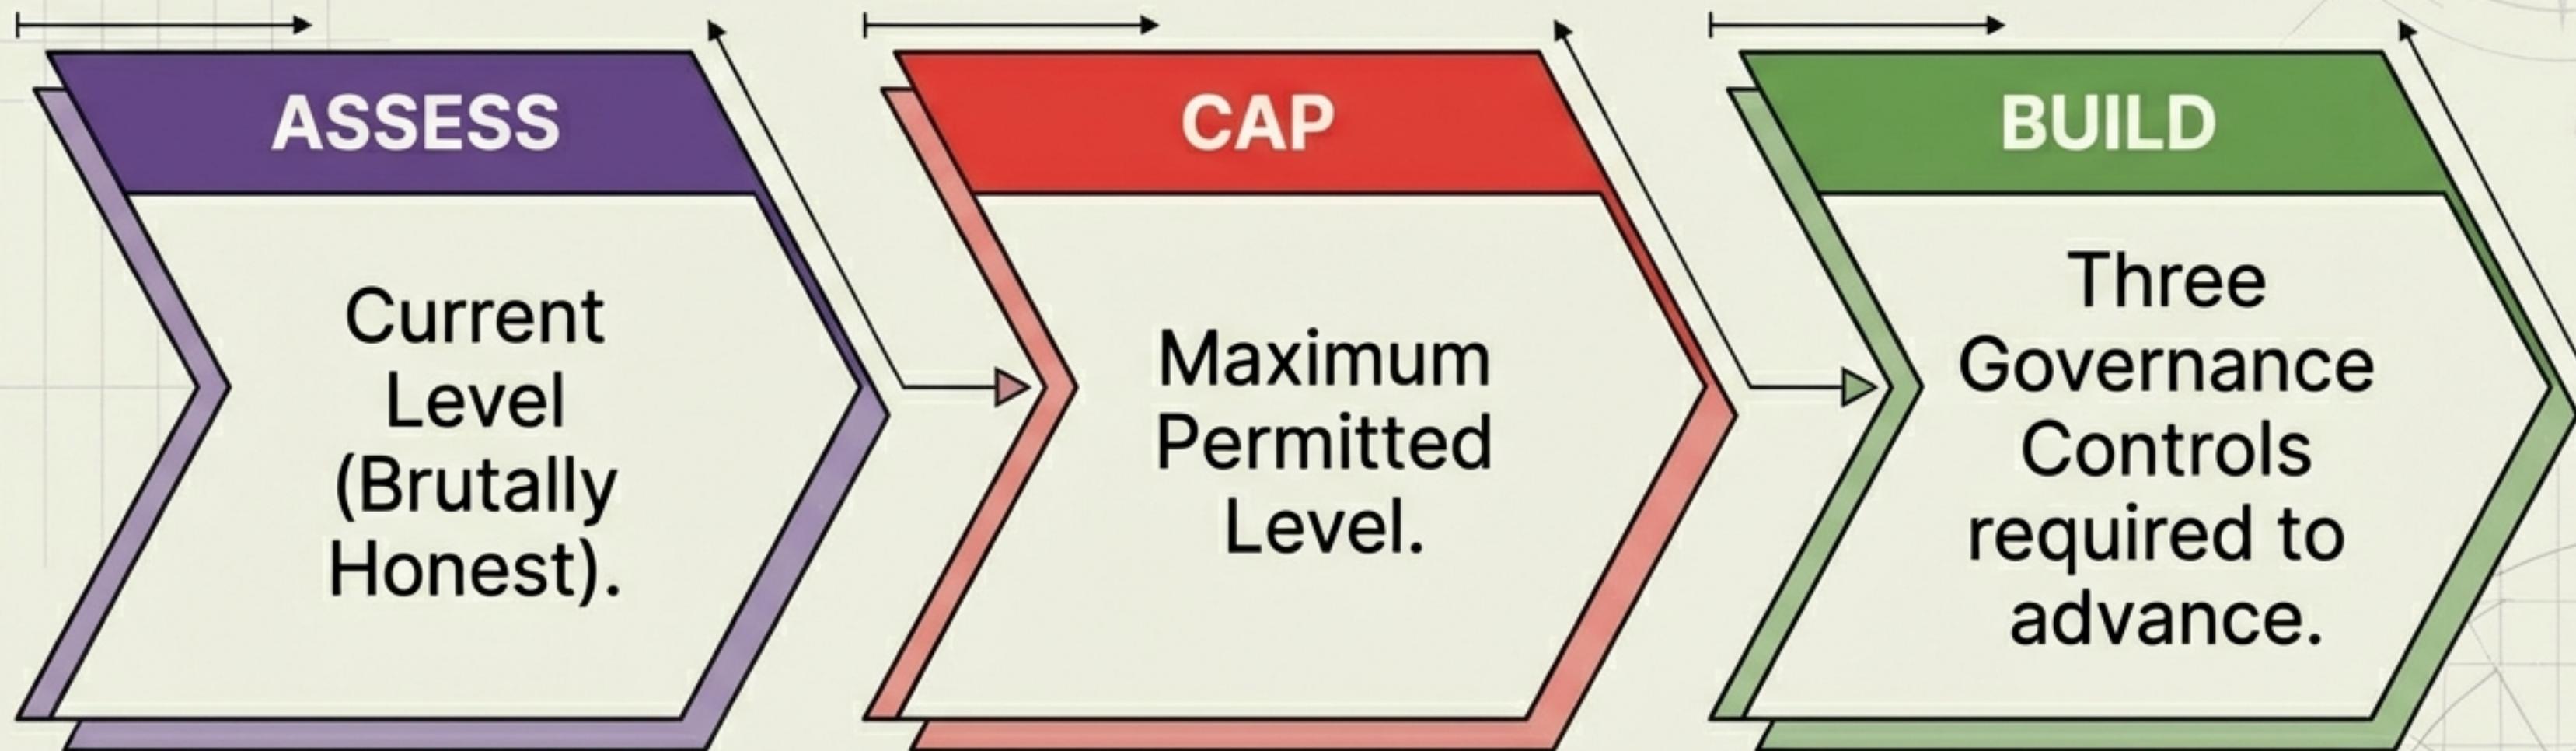
**The Silent Assessment:
Who gets woken up at 3 AM when
the AI fails?**

1. **Which maturity level**, in which domain, with what **failure radius**?
2. Show me one decision that has moved from **human-only** to agent-assisted in the past quarter.
3. What is one **decision category** we've **explicitly** declared non-agentizable?

The gap between the official narrative and the ground truth reveals your actual starting point.

The Implementation Roadmap

Your Next Practical Step



⚠ **Warning:** The most dangerous transition is jumping from Level 1 confidence to Level 3 ambition.

The Board Narrative

Reframing the Investment Case

“Our next investment isn’t in more AI, but in the governance systems that make AI survivable.”

Agentic maturity is not measured by how autonomous our systems are, but by how deliberately we’ve engineered accountability.

Governance IS the product.

The Agentic Maturity Ladder: At a Glance

LEVEL	RISK (COLOR)	STATUS	CONTROL
Level 0	● Risk: Low (Green)	Status: Baseline	SaaS Policy
Level 1	● Risk: Low (Green)	Status: Illusion of Progress	Human Process Owner
Level 2	△ Risk: Medium (Yellow)	Status: Danger Zone	Negative Space / Pre-mortems
Level 3	○ Risk: High (Red)	Complexity Cliff	Circuit Breakers
Level 4	○ Risk: Critical (Purple)	Status: Enterprise Exception	Mathematical Proofs

Internal Reference Document. Use to audit internal proposals.